

# Programming with MPI

## *Introduction*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

April 2010

# Why Use MPI?

CPU<sup>s</sup> got faster at 40% per annum until  $\approx 2003$

Since then, they have got larger but not faster

The number of CPU cores per chip is now increasing

- The solution is to use more CPUs in parallel

MPI (Message Passing Interface) is a tool for that

We will come back to how to obtain MPI later

# Before We Start

You **must** understand your language's **basic features**  
What do you **need**? The course assumes **very little**

- Built-in **datatypes** and **simple calculations**
- The **control statements** and **simple procedures**
- Very simple **1-** and **2-D** arrays
- Very simple use of **I/O** and **text files**

A **very** few exercises need a **little** more

# Course Structure (1)

Start with **essential background** and **basic concepts**  
And running minimal but useful MPI programs

Then move on to facilities used **in practice**  
Based on analysis of initially **twelve** real applications  
Also mention features you might want in the future

Will describe their **underlying concepts** as relevant  
Not well covered in most books and Web pages  
This is helpful for **debugging** and **tuning**

# Course Structure (2)

Also cover **practical** aspects that can cause trouble  
Naturally, based on my personal experience!

Some of these (like **I/O**) are a bit weird  
Will give simple **guidelines** for safe programming

Then give **overview** of more advanced features  
Some are described in books and Web pages  
But **implementations** may not be thoroughly tested

Will **not** go into detail for all of MPI

# Applications

Applications I have looked at include:

Casino, CASTEP, CETEP, CFX11, CPMD,  
CRYSTAL, DLPOLY\_3, Fluent, FFTW,  
mpi\_timer, ONETEP, PARPACK, SPOOLES  
ScaLAPACK and TOMCAT

Only facility course omits entirely is **parallel I/O**

Only in **Fluent** and **DLPOLY\_3** when I looked

Very **specialist** – few people will be interested

# Course Objectives (1)

- The understanding of MPI's **essential concepts**  
How it is likely to be **implemented** (in principle)
- Be able to use **all basic features** of MPI  
For an **empirical** meaning of “**all basic features**”
- Be able to write **highly parallel** HPC code  
Be able to work on almost all **existing** ones
- Be aware of the **ancillary skills** needed

# Course Objectives (2)

- Be able to use **I/O** and other **system interfaces**  
Including knowing something of **what not to do**
- Concepts needed for **debugging** and **tuning**  
Some **experience** of doing so in simple programs
- Knowing what **advanced features** exist in MPI  
So that you don't have to reinvent the wheel
- Also knowing which features are **tricky to use**  
So that you don't use them by accident



# Course Objectives (3)

- This teaches you to program MPI **for real**  
It doesn't skip over anything you **need** to know  
You will still have to look up some **interfaces**  
The intent is that you know **what** to look up
- You will know **why** and **how** things work  
Helps with writing **reliable**, **portable** code  
Minimises confusion when you make a mistake  
And gives a good start with **tuning** your code

All of the above is easier than it looks

# Beyond the Course (1)

Contact your **supervisor** in the first instance  
I am happy for your **supervisor** to contact me

The MPI **standard** home page – final authority  
<http://www.mpi-forum.org/>

Most books / courses skip over **basic concepts**  
And too much time on the more advanced features

This one seems pretty good:  
<http://www.cs.usfca.edu/mpi/>

- This course does **not** follow it!

## Beyond the Course (2)

The materials for this course are available from:

<http://people.ds.cam.ac.uk/nmm1/MPI/>

Several other relevant Computing Service courses  
Some will be mentioned in passing, but see:

<http://people.ds.cam.ac.uk/nmm1/>

They are all “**transferrable skills**” courses  
**Not** part of this MPhil, so get no credit

## Beyond the Course (3)

All of these pages have reliable information  
Most of the Web isn't reliable, of course

[http://www-users.york.ac.uk/~mijp1/teaching/...](http://www-users.york.ac.uk/~mijp1/teaching/.../4th_year_HPC/notes.shtml)  
[.../4th\\_year\\_HPC/notes.shtml](http://www-users.york.ac.uk/~mijp1/teaching/.../4th_year_HPC/notes.shtml)

[http://www.epcc.ed.ac.uk/library/documentation/...](http://www.epcc.ed.ac.uk/library/documentation/.../training/)  
[.../training/](http://www.epcc.ed.ac.uk/library/documentation/.../training/)

<http://www-unix.mcs.anl.gov/mpi/>

# Distributed Memory

One of the basic **parallelism** models

A **program** is run as separate, independent **processes**  
Can be considered as separate **serial** programs

**Distributed memory** means **no** shared data

- The **processes** interact only by **message passing**

May be run on the **same** system or on **separate** ones

# Message Passing

One of the basic **communication** designs

Process **A** sends a message to Process **B**

Process **B** then receives that message

- Think of it as **process-to-process** I/O or Email  
Actually implemented using very similar mechanisms!

Some extra complications, but they use the same idea

# What Is MPI? (1)

- A library callable from Fortran, C (and C++)  
Bindings also available for Python, Java etc.

Primarily for HPC programs on multi-CPU systems  
Assumes a number of processes running in parallel  
Usually with dedicated CPUs (i.e. gang scheduling)

- Essentially all HPC work on clusters uses MPI  
It works nearly as well on multi-core SMP systems
- Poorly for background work (e.g. cycle stealing)

# What Is MPI? (2)

- It is a specialist **communications library**  
Like **POSIX I/O**, **TCP/IP** etc. – but different purpose  
Almost completely **system-independent**
- Using its **interface** is almost never a problem  
If you can use any library, you can use MPI
- Most important step is to understand its **model**  
I.e. the **assumptions** underlying its **design**  
Ditto for **C++**, **POSIX**, **Fortran**, **TCP/IP** and **.NET**



# The MPI Standard (1)

This was a **genuinely open** standardisation process  
Mainly during the second half of the **1990s**

<http://www.mpi-forum.org/docs/docs.html>

**MPI 1** is basic facilities – all most people use  
Most people use only a **small fraction** of it!

**MPI 2** added **extensions** (other facilities)  
Also included the **MPI 1.3** update

**MPI 3** adds yet more – not mentioned in this course

# The MPI Standard (2)

- This is a **standard**, not a **user's guide**  
Designed to be **unambiguous**, not **easy to follow**

As good as **Fortran**, **much** better than **C** or **POSIX**

- But its **order** and **indexing** are ghastly  
⇒ I am still finding new features after a decade
- Use it to look up the precise **specifications**
- Use **something else** to find **what** to look up

# Available Implementations

Two open source versions – MPICH and OpenMPI

You can install as packages or build from source

Most vendors have own, inc. Intel and Microsoft

Usually use shared-memory on multi-core machines

And TCP/IP over Ethernet and other networks

And often InfiniBand on suitable HPC clusters

- But NO code changes are needed!

MPI programs are very portable, and efficiently so

# The MPI Model (1)

You start up **N** independent **processes**

All of them start MPI and use it to communicate

- There is no “**master**” (initial or main process)

Communications may be “**point-to-point**” (pairwise)

- Only two communicating **processes** are involved


Communications may be “**collective**”

All of the **processes** are involved

- They must **all** make the same call, **together**

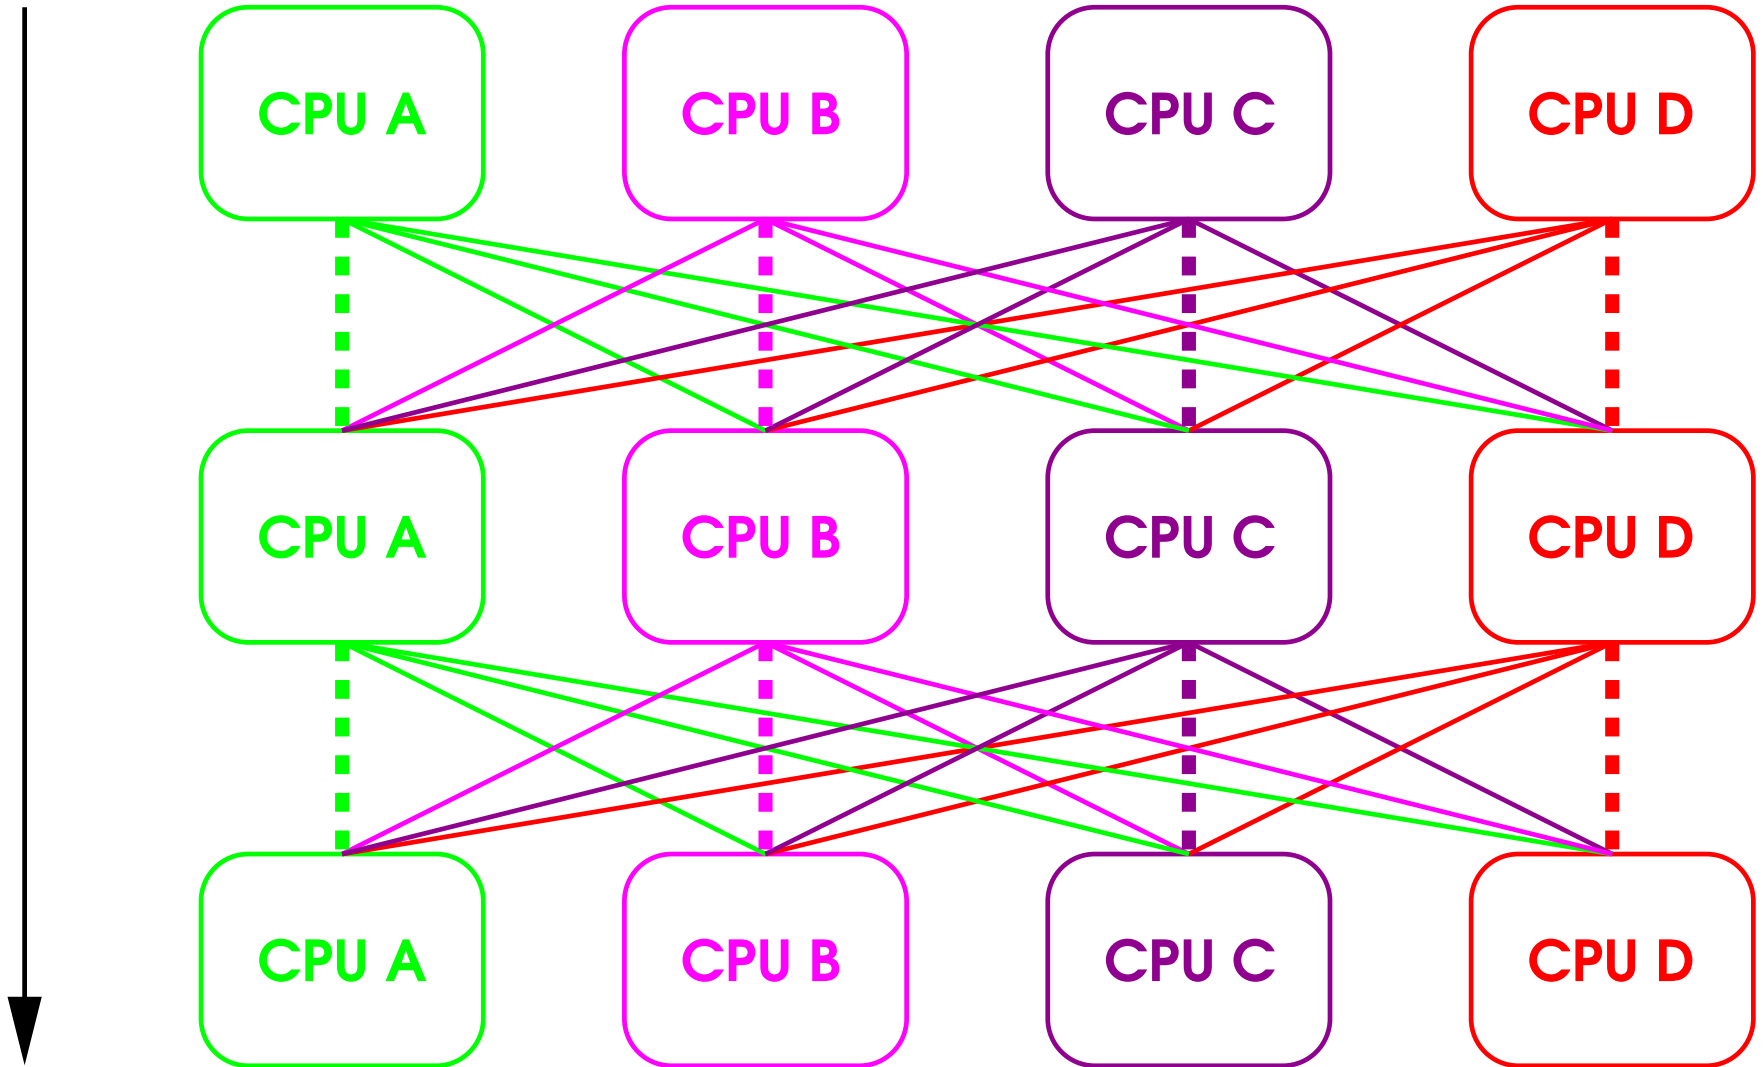
# Point to point Communication

Time



# Collective Communication

Time



# The MPI Model (2)

- **Communication** may not always **synchronise**  
That applies to **collectives** as well as **point-to-point**  
[ The previous picture is misleading in that respect ]
- **Processes** need wait only when they need data  
E.g. a **send** may return before the **receive**  
In theory, this allows for faster execution
- If you want **synchronisation**, you must ask for it  
There are plenty of facilities for doing so

# The MPI Model (3)

Some MPI operations are **non-local**

May involve behind-the-scenes **communication**

Which means they **can hang** if you make an error

And some operations are purely **local**

They can **never hang**, and will return “immediately”

Generally, this matters mainly to MPI **implementors**

- You only need to know that **both forms exist**



# The MPI Model (4)

- Almost everyone uses MPI in **SPMD** mode  
That is **Single Program, Multiple Data**  
You run **N** copies of **one executable**
- The **programs** can execute different **instructions**  
They **don't** have to run in **lockstep** (**SIMD** mode)  
That is **Single Instruction, Multiple Data**  
But start off by designing them to do that
- All **CPUs** are dedicated to your MPI program  
That avoids certain problems I won't describe now

# The MPI Model (5)

SPMD isn't required by MPI, which surprises people  
In theory, don't even need compatible systems  
Could use it on a random collection of workstations

- Don't go there – and not because of MPI

For more detail on the reasons, see:

Parallel Programming: Options and Design

- This course will assume SPMD mode

Many implementations support only SPMD mode

# Communicators


- All communications occur within **communicators**  
A **context**, defining a **group** of **processes**  
Actions in separate **communicators** are independent
- You start with the **communicator** of all processes  
You can **subset** any existing **communicator**

Facilities for that will be described later

- For now, use only **MPI\_COMM\_WORLD**

# Hierarchical Communicators

**MPI\_COMM\_WORLD**



# Number of CPUs (1)

- Parallelism counting is “one, two, many”  
You need to use different algorithms and code

One CPU is necessarily serial programming

Two CPUs are this CPU and the other CPU

Most issues arise only with many CPUs

- Serial codes may not work on many CPUs
- Parallel codes may not work on one CPU
- Two CPU codes may not work on either

## Number of CPUs (2)

MPI **communicators** can have any number of CPUs  
From **zero** CPUs upwards – yes, **no CPUs**

Use **4+** CPUs when debugging generic MPI codes

- Most applications assume at least that many
- This course will cover **only** this case

Otherwise, you need different **code** for:

- **0**: typically do nothing (not executed!)
- **1**: use serial code for this
- **2–3**: a few generic algorithms fail
- **4+**: ‘proper’ parallel working

# Diversion – a Worked Example

Shall now give a worked example of the use of MPI  
Calculate the area of the **Mandelbrot set**

This is to give a feel for what MPI is about  
Don't worry if you don't understand the details  
Every facility used will be explained later

- The whole source is in the extra files  
There are **Fortran 90**, **C** and **C++** versions

# The Mandelbrot Set

This is defined in the complex plane

Consider the recurrence  $x_{n+1} \Leftarrow x_n^2 + c$


With the starting condition  $x_0 = 0$

The **Mandelbrot set** is the set of all  $c$ , such that

$$|x_n| \leq 2, \text{ for all } n$$

This is, er, complicated – let's see a picture





# Calculating its Area

All points within it have  $|c| \leq 2$

It's also symmetric about the X-axis

So we consider just points  $c$ , such that


$$-2 < \operatorname{re}(c) \leq +2$$

$$0 < \operatorname{im}(c) \leq +2$$

Choose a suitable **iteration limit** and **step size**

See if each point stays small for that long

Accumulate the scaled count of those that do



# Example Program

This is the crudest form of numerical integration  
Not strictly **Monte-Carlo**, but is related  
Sometimes a sledgehammer is the best tool!

I have chosen to use **Fortran 90**  
The **C** or **C++** are very similar

Most of it is just the ordinary, serial logic  
I will go through the core of it first

# Testing a Point

```
PURE FUNCTION Kernel (value)
  IMPLICIT NONE
  LOGICAL :: Kernel
  COMPLEX(KIND=DP), INTENT(IN) :: value
  COMPLEX(KIND=DP) :: work
  INTEGER :: n

  work = value
  DO n = 1, maxiters
    work = work**2 + value
    IF (ABS(REAL(work)) > 2.0 .OR.      &
        ABS(AIMAG(work)) > 2.0) EXIT
  END DO
  Kernel = (ABS(WORK) <= 2.0)
END FUNCTION Kernel
```

# Scanning an Area

```
PURE FUNCTION Shell (lower, upper)
  IMPLICIT NONE
  REAL(KIND=DP) :: Shell
  COMPLEX(KIND=DP), INTENT(IN) :: lower, upper
  COMPLEX(KIND=DP) :: work

  Shell = 0.0_DP
  work = CMPLX(REAL(lower),      &
               AIMAG(lower)+step/2.0_DP,KIND=DP)
  DO WHILE (AIMAG(work) < AIMAG(upper))
    DO WHILE (REAL(work) < REAL(upper))
      IF (Kernel(work)) Shell = Shell+step**2
      work = work+step
    END DO
    work = CMPLX(REAL(lower),AIMAG(work)+step,KIND=DP)
  END DO
END FUNCTION Shell
```

# MPI Initialisation

```
LOGICAL, PARAMETER :: UseMPI = .True.
```

```
INTEGER, PARAMETER :: root = 0
```

```
INTEGER :: maxiters, error, nprocs, myrank
```

```
REAL(KIND=DP) :: buffer_1(2), step, x
```

```
IF (UseMPI) THEN
```

```
    CALL MPI_Init(error)
```

```
    CALL MPI_Comm_size(MPI_COMM_WORLD,nprocs,error)
```

```
    CALL MPI_Comm_rank(MPI_COMM_WORLD,myrank,error)
```

```
ELSE
```

```
    nprocs = 1
```

```
    myrank = root
```

```
END IF
```

# Divide Area into Domains

```
COMPLEX(KIND=DP), ALLOCATABLE :: buffer_2(:,,:)
```

```
IF (myrank == root) THEN
```

```
    ALLOCATE(buffer_2(2,nprocs))
```

```
    buffer_2(1,1) = CMPLX(-2.0_DP,0.0_DP,KIND=DP)
```

```
    DO i = 1,nprocs-1
```

```
        x = i*2.0_DP/nprocs
```

```
        buffer_2(2,i) = CMPLX(2.0_DP,x,KIND=DP)
```

```
        buffer_2(1,i+1) = CMPLX(-2.0_DP,x,KIND=DP)
```

```
    END DO
```

```
    buffer_2(2,nprocs) = CMPLX(2.0_DP,2.0_DP,KIND=DP)
```

```
ELSE
```

```
    ALLOCATE(buffer_2(2,1))    ! This is not actually used
```

```
END IF
```



# Reading the Parameters

```
INTEGER :: maxiters  
REAL(KIND=DP) :: step
```

```
IF (myrank == root) THEN  
    READ *, maxiters, step  
    IF (maxiters < 10) THEN  
        PRINT *, 'Invalid value of MAXITERS'  
        CALL MPI_Abort(MPI_COMM_WORLD,1,error)  
    END IF  
    IF (step < 10.0_DP*EPSILON(step) .OR. step > 0.1_DP) THEN  
        PRINT *, 'Invalid value of STEP'  
        CALL MPI_Abort(MPI_COMM_WORLD,1,error)  
    END IF  
END IF
```

# Distribute the Data (1)

```
REAL(KIND=DP) :: buffer_1(2)
COMPLEX(KIND=DP), ALLOCATABLE :: buffer_2(:, :)
COMPLEX(KIND=DP) :: buffer_3(2)

IF (myrank == root) THEN
    buffer_1(1) = maxiters
    buffer_1(2) = step
END IF
```

# Distribute the Data (2)

```
IF (UseMPI) THEN
    CALL MPI_Bcast(      &
        buffer_1,2,MPI_DOUBLE_PRECISION,      &
        root,MPI_COMM_WORLD,error)
    maxiters = buffer_1(1)
    step = buffer_1(2)
    CALL MPI_Scatter(      &
        buffer_2,2,MPI_DOUBLE_COMPLEX,      &
        buffer_3,2,MPI_DOUBLE_COMPLEX,      &
        root,MPI_COMM_WORLD,error)
ELSE
    buffer_3 = buffer_2(:,1)
END IF
```

# Accumulate in Parallel

```
buffer_1(1) = Shell(buffer_3(1),buffer_3(2))
IF (UseMPI) THEN
    CALL MPI_Reduce(      &
        buffer_1(1),buffer_1(2),      &
        1,MPI_DOUBLE_PRECISION,      &
        MPI_SUM,root,MPI_COMM_WORLD,error)
ELSE
    buffer_1(2) = buffer_1(1)
END IF
```

# Print Results and Terminate

```
IF (myrank == root) THEN
    PRINT '(A,F6.3)',      &
        'Area of Mandelbrot set is about',      &
        2.0_DP*buffer_1(2)
END IF
IF (UseMPI) THEN
    CALL MPI_Finalize(error)
END IF
```

# So What Happens?

Running with parameters '10000 0.001'

We get about 1.508 (true result is about 1.506)

<u>Number of processors</u>	<u>Elapsed time taken</u>
1	67 seconds
4	46 seconds
16	23 seconds

Not very **scalable**, is it? That is quite common  
**Using** MPI is much easier than **tuning** it

# Doing Better (1)

There is an alternative Fortran 90 version, too  
Generates all of the points and randomises them  
Each processor has a roughly matching workload

It is a store hog, and takes some time to start

<u>Number of processors</u>	<u>Elapsed time taken</u>
1	70 seconds
4	19 seconds
16	8 seconds

It would scale better with more points

## Doing Better (2)

There is a better way than even that, too  
Covered in the **Problem Decomposition** lecture  
The first practical of that gets you to do it

Suitable for **embarrassingly parallel** problems  
E.g. **parameter searching** and **Monte-Carlo** work  
**Mandelbrot set** was merely a convenient example

But that's a lot later . . .