

Programming with MPI

Using MPI

Nick Maclaren

nmm1@cam.ac.uk

May 2008

Warning

This lecture covers a huge number of minor points
Including **all** of the ‘**housekeeping**’ facilities

Don't try to remember all details, initially

- Try to remember **which facilities** are included

Refer back to this when doing the practicals

It's a **lot** easier than it looks at first

Using MPI

- By default, all actual **errors** are fatal
MPI will produce some kind of an **error message**
With luck, the **whole program** will then stop

Can ask to do your own **error handling** – see later

- Use **one** interface: **Fortran** or **C**
C++ can use **C**, possibly with '**extern "C"**'

Yes, you can mix them – but it's advanced use

Function Declarations

There are **proformas** for all functions used
Anything merely mentioned is omitted, for clarity

Interfaces/Fortran

Interfaces/C

The examples don't give the **syntax** in detail
Check those files when doing the practicals

Warning: they do not match the **current MPI**
Use **INTENT(IN)** and **const**, as in the latest one
Will be **no problem** if you just **use the functions**

MPI's Fortran Interface (1)

- If possible, include the statement: `USE mpi`
- If not, use: `INCLUDE 'mpif.h'`
after all “`USE`”s and “`IMPLICIT`”

Note the first is “`mpi`” and the second “`mpif.h`”
If both fail, usually a usage / installation problem

All MPI names start with `MPI_`

- Don't declare `names` starting `MPI_` or `PMPI_`
Names `PMPI_` are used for `profiling`

MPI's Fortran Interface (2)

Boolean values (true/false) are LOGICAL

Process numbers, error codes etc. are INTEGER

Element counts etc. are also plain INTEGER

This isn't a problem on any current system

Almost all MPI constants are Fortran constants

The only exception in this course is MPI_IN_PLACE

Arrays start at one, where it matters

MPI's Fortran Interface (3)

Type-generic arguments are a kludge
MPI relies on Fortran not noticing them
Will describe the issues later

MPI 3 and Fortran TS 29113 fixes them properly

For now, just pass arrays of any type
If the compiler objects, ask for help
Some guidelines on how in a later lecture

MPI's Fortran Interface (4)

Handles (e.g. **communicators**) are **opaque** types

[One you can't break apart and look inside]

Undocumented and unpredictable **INTEGER** values

Use built-in **equality comparison** and **assignment**

Call MPI functions for **all** other operations

I.e. MPI returns **INTEGER** values as **tokens**

If their **values** match, they are the same **token**

MPI's Fortran Interface (5)

- Almost all MPI functions are **SUBROUTINES**
The **final argument** returns an **INTEGER** error code

Success returns **MPI_SUCCESS** (always zero)

Failure codes are **implementation dependent**

Only a **very few** exceptions: mainly **MPI_Wtime**

All **results** are returned through **arguments**

MPI's Fortran Interface (6)

As people will know, default **REAL** is a disaster
DOUBLE PRECISION is tedious and out-of-date

Start all **procedures**, **modules** etc. with

USE double

USE mpi

IMPLICIT NONE

There is a suitable file **Programs/double.f90**

Ask for help if you don't know how to use it

MPI's C Interface (1)

- C++ uses the C interface

So C++ people need to listen to this section

Include the statement: `#include "mpi.h"`

If that doesn't work, see the next slide

All MPI names start with `MPI_`

- Don't declare names starting `MPI_` or `PMPI_`

Names `PMPI_` are used for profiling

C++: Using the C Interface

If simple `#include "mpi.h"` doesn't work, try this:

```
extern "C" {  
    #include "mpi.h"  
}
```

Usually, one or the other approach will work

Or put your MPI into a separate file called `*.c`

Write it in pure C, and use a C compiler on it

Now use `extern "C"` to use that from your C++

MPI's C Interface (2)

Boolean values (true/false) are `int`, as usual

Process numbers, error codes etc. are `int`

Element counts etc. are also plain `int`

This isn't a problem on any current system

Type-generic arguments are `void *`

These are called “choice” arguments by MPI

MPI's C Interface (3)

Almost all MPI constants are C initialization constants
NOT usually preprocessor or integer constants

- Cannot use in case, array sizes etc.

Only maximum sizes are preprocessor constants

Arrays start at zero, where it matters

MPI's C Interface (1)

Handles (e.g. communicators) are opaque types

Names are set up by typedef and are scalars

Use built-in equality comparison and assignment

Call MPI functions for all other operations

The main such opaque types are:

MPI_Comm, MPI_Datatype, MPI_Errhandler,
MPI_Group, MPI_Op, MPI_Request,
MPI_Status

MPI's C Interface (2)

- Almost all MPI functions return an **error code**

This is the **function result** as an **int**

Can ignore it, if using **default** error handling

Success returns **MPI_SUCCESS** (must be zero)

Failure codes are **implementation dependent**

Only a **very few** exceptions: mainly **MPI_Wtime**

- All **results** are returned through **arguments**

MPI and C++

MPI 2.0 introduced a C++ interface in 1997

It's significantly better in a great many respects

However, MPI 2.2 deprecated it in 2009

Its recommendation is to use the C interface

- MPI 3.0 has deleted it
 - This course will teach only the C interface
- Handouts and materials cover the C++ one, too
- When calling C from C++, use the C examples

More on Interfaces

- That is all you need for now

We will return to language interfaces later

- Advanced language facilities to avoid
- Interfaces for advanced MPI programming
- Performance and optimisation issues

Starting and Stopping

- For now, we will ignore error handling

All processes must start by calling `MPI_Init`

And, normally, finish by calling `MPI_Finalize`

- These are effectively **collectives**

Call both at predictable times, or risk confusion

- You **can't** restart MPI after `MPI_Finalize`

`MPI_Init` must be called **exactly** once

Fortran Startup/Stopping

Fortran argument decoding is behind the scenes

```
USE double
USE mpi
IMPLICIT NONE
INTEGER :: error

CALL MPI_Init ( error )
CALL MPI_Finalize ( error )
END
```

If that doesn't work, see the MPI documentation

- Though you will probably need to ask for help

C Startup/Stopping

MPI_Init takes the **addresses** of **main**'s arguments

- You **must** call it before **decoding** them

Some **implementations** change them in **MPI_Init**

```
#include "mpi.h"
```

```
int main (int argc , char * argv [] ) {  
    MPI_Init ( & argc , & argv ) ;  
    MPI_Finalize ( ) ;  
    return 0 ;  
}
```

Aside: Examples

I will omit the following statements, for brevity:

```
USE double  
USE mpi  
IMPLICIT NONE
```

```
#include "mpi.h"
```

Include them in any “module” where you use MPI

Don't rely on implicit declaration

Version Numbers

MPI 1.2 and up provide version number information

- Not needed for simple use, as in this course
All versions of MPI are essentially compatible

Constants MPI_VERSION, MPI_SUBVERSION

Set to 1, 3 for MPI 1.3 or 2, 2 for current MPI 2

There is also a function MPI_Get_version

Which can be called even before MPI_Init

Testing MPI's State (1)

You can test the state of MPI on a process

- This is needed only when writing library code

Fortran example:

```
LOGICAL :: started , stopped
```

```
INTEGER :: error
```

```
CALL MPI_Initialized ( started , error )
```

```
CALL MPI_Finalized ( stopped , error )
```


Testing MPI's State (2)

C example:

```
int started , stopped , error ;
```

```
error = MPI_Initialized ( & started ) ;
```

```
error = MPI_Finalized ( & stopped ) ;
```

Global Communicator

The global **communicator** is predefined:

MPI_COMM_WORLD

It includes all usable **processes**

e.g. the **<n>** set up by “**mpiexec -n <n>**”

Many applications use only this **communicator**

- Almost all of this course does, too

There is one lecture on **communicators**

Process Rank

The **rank** is the **process**'s index
always within the context of a **communicator**

A **rank** is an integer from 0 to $\langle n \rangle - 1$

Yes, this applies to **Fortran**, too

There is one predefined **rank** constant:

MPI_PROC_NULL – no such **process**

Don't assume this is negative – or that it isn't

We shall describe the use of it when relevant

Information Calls (1)

`MPI_Comm_size` returns the number of processes

`MPI_Comm_rank` returns the local process number

Fortran example:

```
INTEGER :: nprocs , myrank , error
CALL MPI_Comm_size (      &
    MPI_COMM_WORLD , nprocs , error )
CALL MPI_Comm_rank (      &
    MPI_COMM_WORLD , myrank , error )
```

Remember `&` means continuation in Fortran

Information Calls (2)

C example:

```
int nprocs , myrank , error ;
```

```
error = MPI_Comm_size ( MPI_COMM_WORLD ,  
                        & nprocs ) ;
```

```
error = MPI_Comm_rank ( MPI_COMM_WORLD ,  
                        & myrank ) ;
```

Information Calls (3)

You can query the local **processor name**

A **string** of length **MPI_MAX_PROCESSOR_NAME**

Fortran example:

```
CHARACTER ( LEN =      &  
            MPI_MAX_PROCESSOR_NAME ) :: procname  
INTEGER :: namelen , error  
  
CALL MPI_Get_processor_name ( procname ,      &  
                               namelen , error )
```

Information Calls (4)

C example:

```
char procname [ MPI_MAX_PROCESSOR_NAME + 1 ];  
int namelen , error ;  
  
error = MPI_Get_processor_name ( procname ,  
                                & namelen ) ;  
procname [ namelen ] = '\0';
```

Information Calls (5)

MPI_Wtime gives **elapsed time** (“wall-clock time”)
Seconds since an unspecified **starting point**

The **starting point** is fixed for a **process**
Doesn't change while the **process** is running

I have seen **start of process**, **system boot time**,
Unix epoch and **00:00 Jan. 1st 1900**

MPI_Wtick similar but gives timer **resolution**
Few people bother – but it's there if you want it

Information Calls (6)

Fortran:

```
REAL(KIND=KIND(0.0D0)) :: now  
now = MPI_Wtime ( )
```

C:

```
double now ;  
now = MPI_Wtime ( ) ;
```

Information Calls (7)

Anywhere from `MPI_Init` to `MPI_Finalize`

They are all purely **local** operations

Use them as often as you need them

`MPI_Comm_size` same result on all **processes**

- Others may give **different** ones on each **process**
- That includes `MPI_Wtime`'s starting point
As well as the value returned from `MPI_Wtick`

Barrier Synchronisation (1)

MPI_Barrier synchronises all **processes**

They all wait until they have all entered the call

Then they all start up again, independently

- The **only** collective that **synchronises**

We will come back to this later

Barrier Synchronisation (2)

Fortran example:

```
INTEGER :: error
```

```
CALL MPI_Barrier ( MPI_COMM_WORLD , error )
```

C example:

```
int error ;
```

```
error = MPI_Barrier ( MPI_COMM_WORLD ) ;
```

Abandoning All Hope (1)

`MPI_Abort` is the emergency stop

- Always call it on `MPI_COMM_WORLD`

Not a `collective` but should stop all `processes`
and, on `most` systems, it `usually` does ...

- Outstanding `file output` is often lost

Far better to stop normally, if at all possible

I.e. all `processes` call `MPI_Finalize` and exit

- `MPI_Abort` is the `emergency` stop

Abandoning All Hope (2)

Fortran:

```
INTEGER :: error
```

```
CALL MPI_Abort ( MPI_COMM_WORLD ,      &  
                <failure code> , error )
```

C:

```
int error ;
```

```
error = MPI_Abort ( MPI_COMM_WORLD ,  
                  <failure code> ) ;
```

Lookahead to I/O

I/O in parallel programs is **always** tricky
It's worse in MPI, because of MPI's portability
Each **type of parallelism** has different oddities

- For now, just write to **stdout** or **stderr**
And the default output in **Fortran**, of course
It will work well enough for the examples
Lines may be **interleaved** with each other
- We will come back to using I/O later

First Practical

You can actually do quite a lot with just these

Start by writing a trivial test program

Then writing a **command spawner**

This is very useful, and there are several around

- Yes, some practical uses **ARE** that simple!

You may not know enough **C/Fortran** – if so, skip it

Use any language you like, that can call MPI

Examples will be in **Fortran** and **C**

Compiling and Running

This is all very **implementation-dependent**, of course
But, on most systems, do something like this:

Compile and link using **mpif90**, **mpicc**, **mpiCC**

Run using “**mpiexec -n <n> <program> [args ...]**”

<n> is the number of **processes** to use

When using a **job scheduler** (**queuing system**)
you may need to put the latter in a script

- This course will use MPI only in **SPMD** mode

PWF/MCS/DS Linux Usage

The PWF/MCS/DS uses mainly dual core systems
All the examples will work, but with odd timings

This course teaches core-independent MPI use
This is just another aspect of portability

It uses gfortran, gcc and OpenMPI
I recommend using -Wall -Wextra -g -O3
If so, ignore a few warnings – but only those

Ignorable Warnings

Fortran:

Warning: Procedure '...' called with an implicit interface at (1)

For most of the **MPI** calls – but only those

C/C++:

/usr/local/OPENMPI/include/mpi.h:220: warning:
ISO C90 does not support 'long long'

C++:

/usr/local/OPENMPI/include/openmpi/ompi/mpi/cxx/comm_inln.h:...
warning: unused parameter '...'

Regrettably, there are quite a lot of these

Instructions

If running **Microsoft Windows**, **CTRL-ALT-DEL**
Select **Restart** and then **Linux**
Log into **Linux** and start a shell and an editor
Create programs called **prog.f90**, **prog.c**, **prog.cpp**.

- Run by typing commands like
mpif90 prog.f90, **mpicc prog.c**, **mpiCC prog.cpp**
mpiexec -n 4 a.out
- Analyse what went wrong
- Fix bugs and retry