

Programming with MPI

Datatypes and Collectives

Nick Maclaren

nmm1@cam.ac.uk

May 2008

Transfer Procedures

These need to specify one or more **transfer buffers**
Used to **send** or **receive** data, or both

These are specified using three **arguments**:

- The **address** of the buffer

- The **size** of the buffer

- The base **datatype** of the buffer

They also need to specify some **control information**

- The **root process** for **1:all** transfers

- The **communicator** to be used for the collective

Transfer Buffers (1)

MPI **transfers** use **vectors** (i.e. **1-D** arrays)

The base element **datatypes** are always **scalars**

They all include an **element count** argument
i.e. the **length** of the **vector** in **elements**

- The **arguments** are **type-generic** (**choice**)

Declared as “**void ***” in **C**

Fortran relies on no checking (see later)

- The **datatype** is passed as a separate **argument**

Transfer Buffers (2)

The **vectors** are always **contiguous arrays**
Each **element** immediately follows its **predecessor**

Like **Fortran 77** or **C arrays**, not all of **Fortran 90**
Return to **Fortran 90 assumed shape** arrays later

For example, consider transferring **100 integers**
The **element count** is **100**

These are declared like:

Fortran: **INTEGER BUFFER (100)**

C: **int buffer [100] ;**

C++ Classes

But what about C++ library containers?

In general, they are not contiguous arrays

- But `<vector>`, `<array>` and `<string>` are
⇒ Though not `<vector<bool> >`!

In all cases, `& object . front ()` is the data address

So is `& object [0]`, and many people use that instead

Conversion to `void *` in the call is automatic

- This applies far more generally than just to MPI
Also, similar remarks apply to libraries like Boost

Datatypes (1)

Datatypes are MPI constants, not **language types**

There is a fairly complete set that are **built-in**

- Note that does **NOT** mean **language constants**

Each **datatype** has an associated **size**

- **Count** and **offsets** are in units of that

Exactly as in **Fortran** or **C arrays**

```
double buffer [ 100 ] ;  
MPI_Bcast ( buffer , 100 , MPI_DOUBLE ,  
           root , MPI_COMM_WORLD )
```

Datatypes (2)

The MPI and language **datatypes** must match
Some exceptions, but I suggest avoiding them

- You will **not** get warned if you make an error

As in **K&R C**, **C casts** and **Fortran 77**

There is no **C++** or **Fortran 90** type-checking

In theory, a compiler could detect a mismatch

But it would have to be “**MPI aware**” and none are

Datatypes (3)

Here is a **sample** of recommended datatypes
All that you need for the first examples
We will come back to these in more detail later

Fortran:

MPI_INTEGER

MPI_DOUBLE_PRECISION

C:

MPI_INT

MPI_DOUBLE

Collectives (1)

We have already used `MPI_Barrier`

All of the others involve some `data transfer`

- All `processes` in a `communicator` are involved
For use on a `subset`, create another `communicator`
We shall come back to that later

- All `datatypes` and `counts` must be the same
A few, obscure exceptions – not recommended
Obviously the `communicator` must be, too

Collectives (2)

- All of the **buffer addresses** may be different
MPI **processes** don't share any **addressing**

This generalises in more advanced use

The **data layout** may be different – see later

- Match the **communicator**, **datatypes** and **counts**
And call all of the **collectives** “at the same time”

- Easiest to achieve using the **SPMD** model
You can code just one **collective call**

Collectives (3)

Some **collectives** are **asymmetric** (**1:all**)

E.g. broadcast from one proc. to all **communicator**

That means **all processes** – including itself

Those all have a **root process** argument

This also **must** be the same on all **processes**

Any **process** can be specified – not just **zero**

Symmetric ones don't have that argument

For example, **MPI_Barrier** doesn't

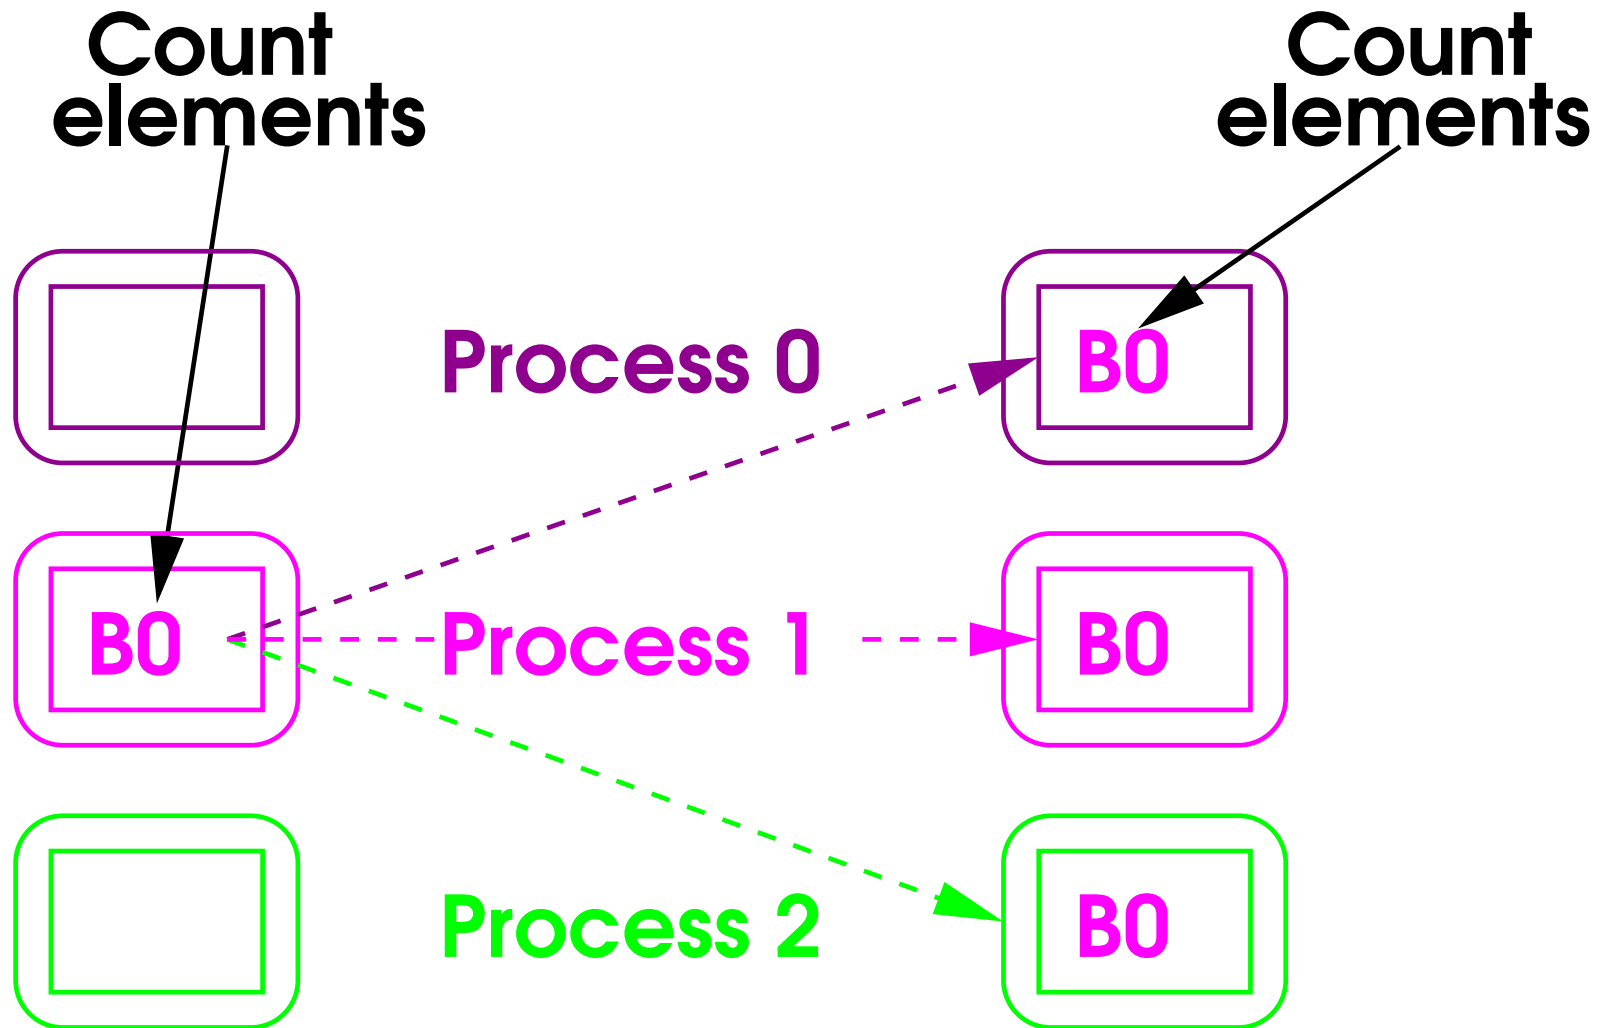
Collectives (4)

- Most use separate **send** and **receive** buffers
Both for **flexibility** and for **standards conformance**
- Usually specify the **datatype** and **count** for each
Needed for advanced features not covered here

MPI uses only the **arguments** it needs
I.e. unused ones are completely ignored

- Set them all compatibly – it is much safer!
Keep all **datatypes** and **counts** the same

Broadcast



Broadcast (1)

Broadcast copies the same data from the **root** to all **processes** in the **communicator**

Fortran example:

```
REAL(KIND=KIND(0.0D0)) :: buffer ( 100 )  
INTEGER , PARAMETER :: root = 3  
INTEGER :: error  
CALL MPI_Bcast ( buffer , 100 ,      &  
                MPI_DOUBLE_PRECISION , root ,      &  
                MPI_COMM_WORLD , error )
```

Broadcast (2)

C example:

```
double buffer [ 100 ] ;  
int root = 3 , error ;  
error = MPI_Bcast ( buffer , 100 , MPI_DOUBLE ,  
    root , MPI_COMM_WORLD ) ;
```

C++ using C interface example:

```
vector<double> buffer ( 100 ) ;  
int root = 3 ;  
error = MPI_Bcast ( & buffer . front ( ) , 100 ,  
    MPI_DOUBLE , root , MPI_COMM_WORLD ) ;
```

Multiple Transfer Buffers

Many **collectives** need one **buffer** per **process**

For example, take a **1 \Rightarrow N** **scatter** operation

The **root** sends different data to each **process**

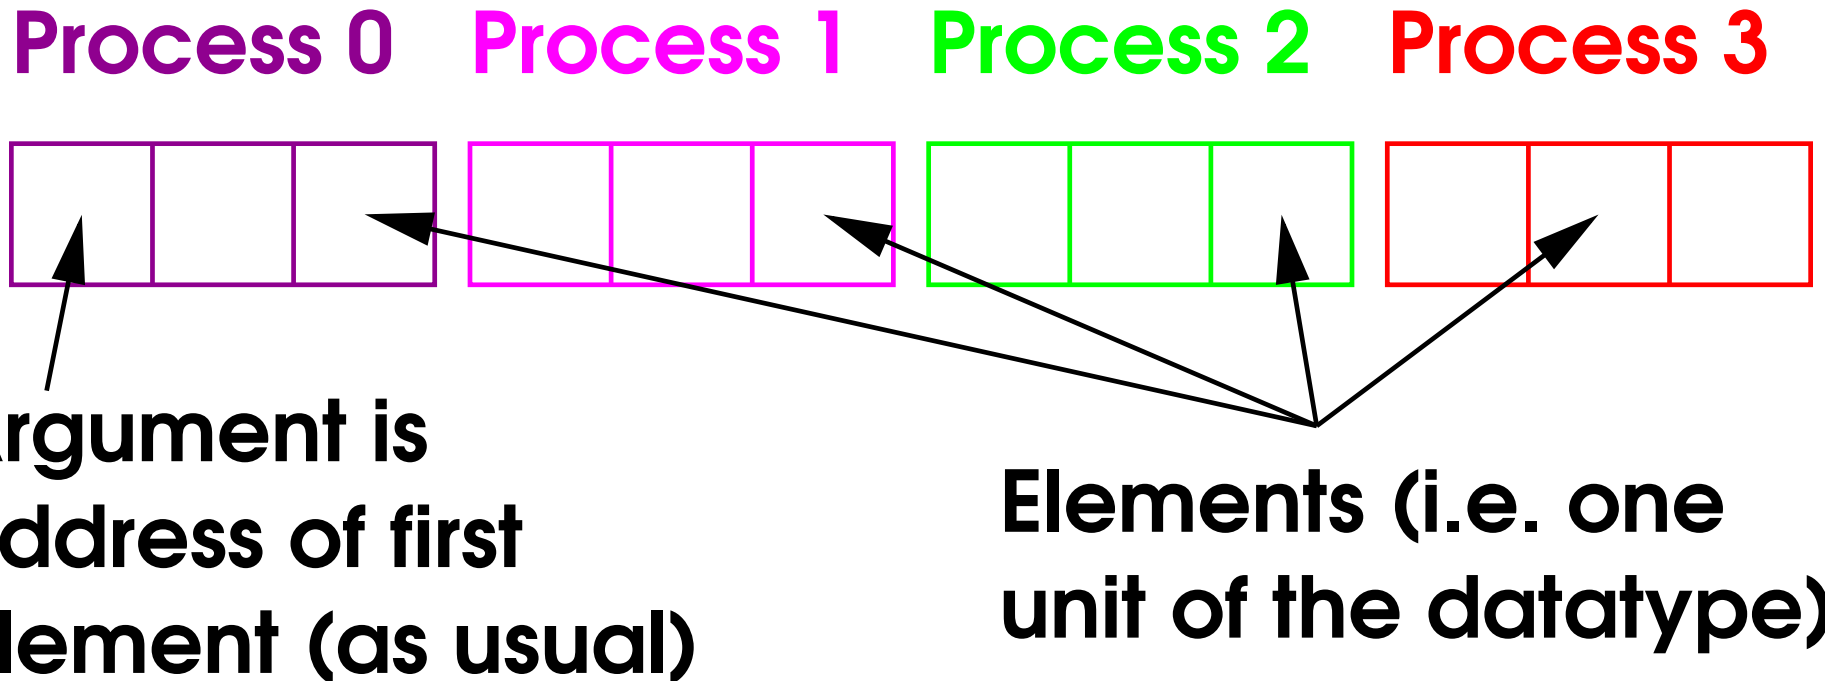
Each pairwise **transfer buffer** is concatenated
in the order of **process numbers** (i.e. **0...N-1**)

Size of source = **N** * **size of each result**

Multiple Transfer Buffers

This is for 4 processes

A count (vector length) of 3



Size Specifications (1)

Size specifications are slightly counter-intuitive
That is done for consistency and simplicity

You specify the size of each pairwise transfer
MPI will deduce the total size of the buffers
I.e. it will multiply by process count, if needed

- The process count is implicit
It is taken from the communicator
I.e. the result from `MPI_Comm_size`

Size Specifications (2)

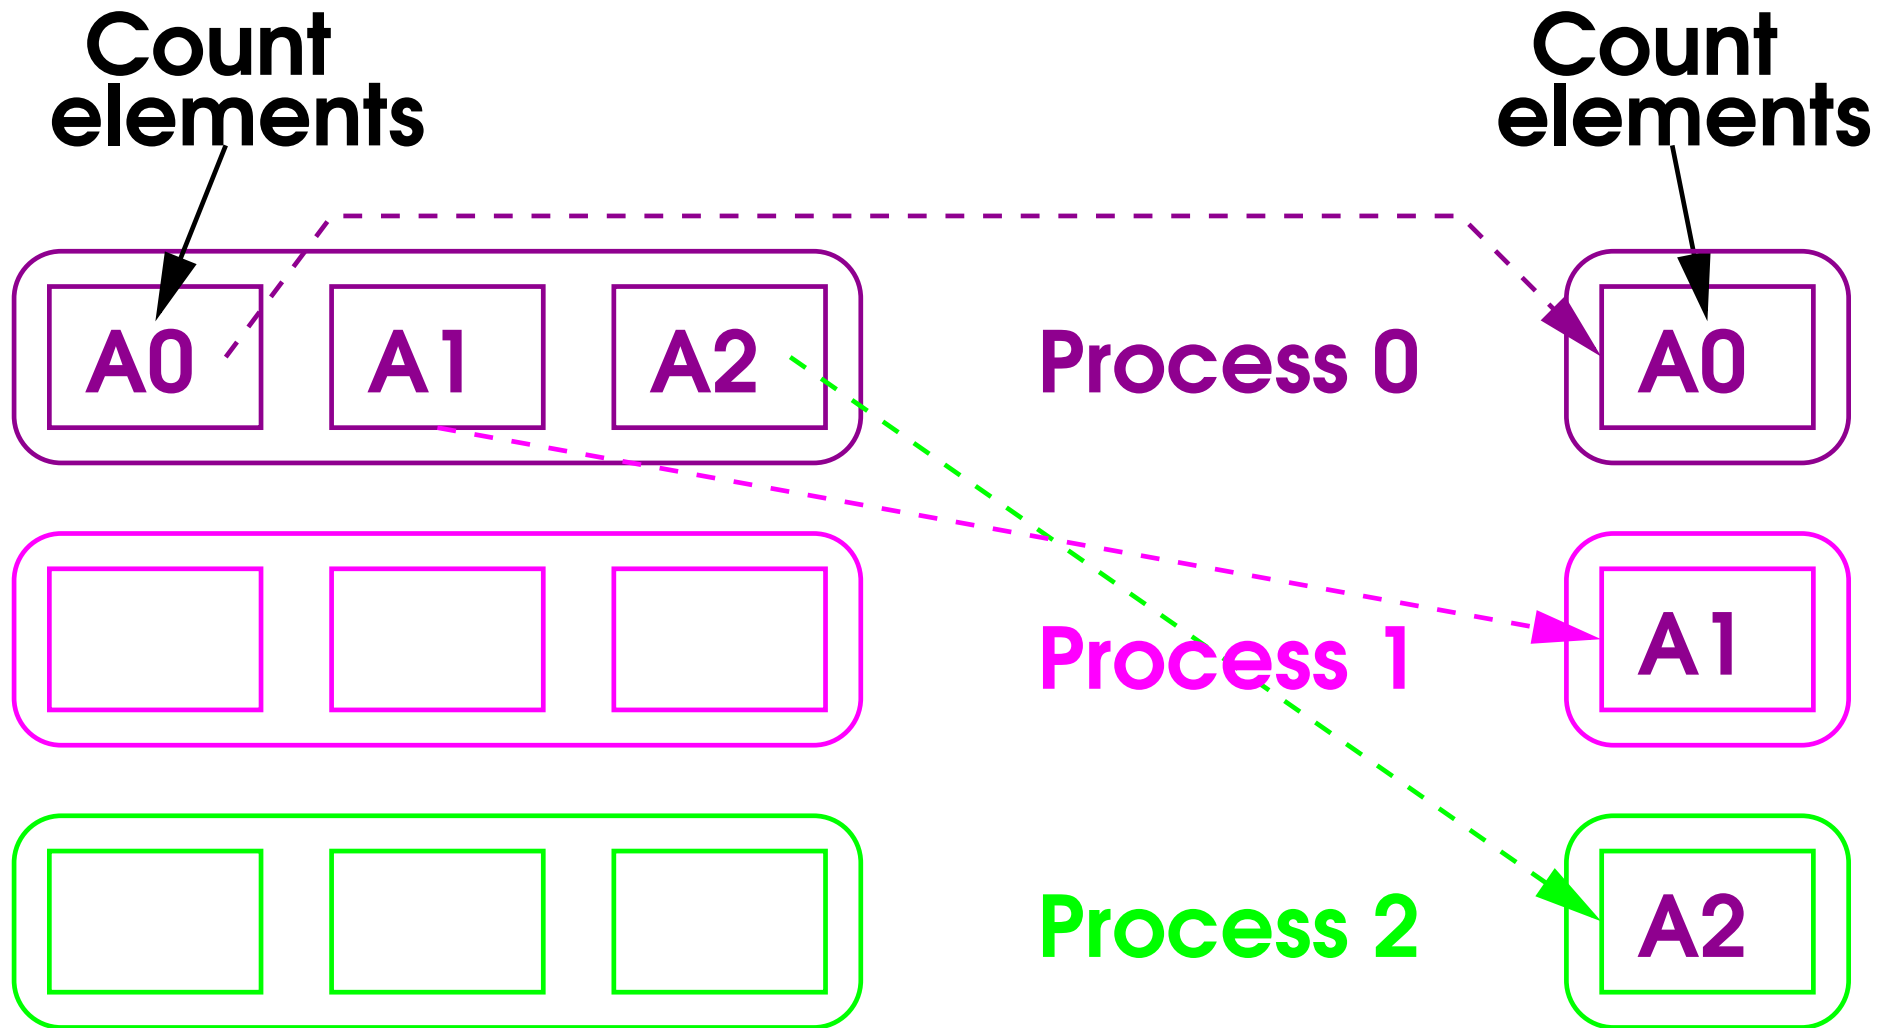
“**void ***” defines no **length** in **C**

Nor does “**<type> :: buffer(*)**” in **Fortran**

- It is up to you to get it right
- No compiler can trap an error with that

We shall use **scatter** as our first example
This is one **process** sending different data
to every process in the **communicator**

Scatter



Scatter (1)

Scatter copies different data from the **root** to all **processes** in the **communicator**

The **send buffer** is used only on the **root**

The **receive buffer** is used on all **processes**

Following examples assume ≤ 30 **processes**

Specified **only** in the **send buffer** size

- Note the differences in the **buffer declarations**

Scatter (2)

Fortran example:

```
REAL(KIND=KIND(0.0D0)) ::      &  
    sendbuf ( 100 , 30 ) , recvbuf ( 100 )  
INTEGER , PARAMETER :: root = 3  
INTEGER :: error  
CALL MPI_Scatter (      &  
    sendbuf , 100 , MPI_DOUBLE_PRECISION ,      &  
    recvbuf , 100 , MPI_DOUBLE_PRECISION ,      &  
    root , MPI_COMM_WORLD , error )
```

Scatter (3)

C example:

```
double sendbuf [ 30 ] [ 100 ] , recvbuf [ 100 ] ;  
int root = 3 , error ;  
error = MPI_Scatter (   
    sendbuf , 100 , MPI_DOUBLE ,  
    recvbuf , 100 , MPI_DOUBLE ,  
    root , MPI_COMM_WORLD )
```

Scatter (4)

C++ using C interface example:

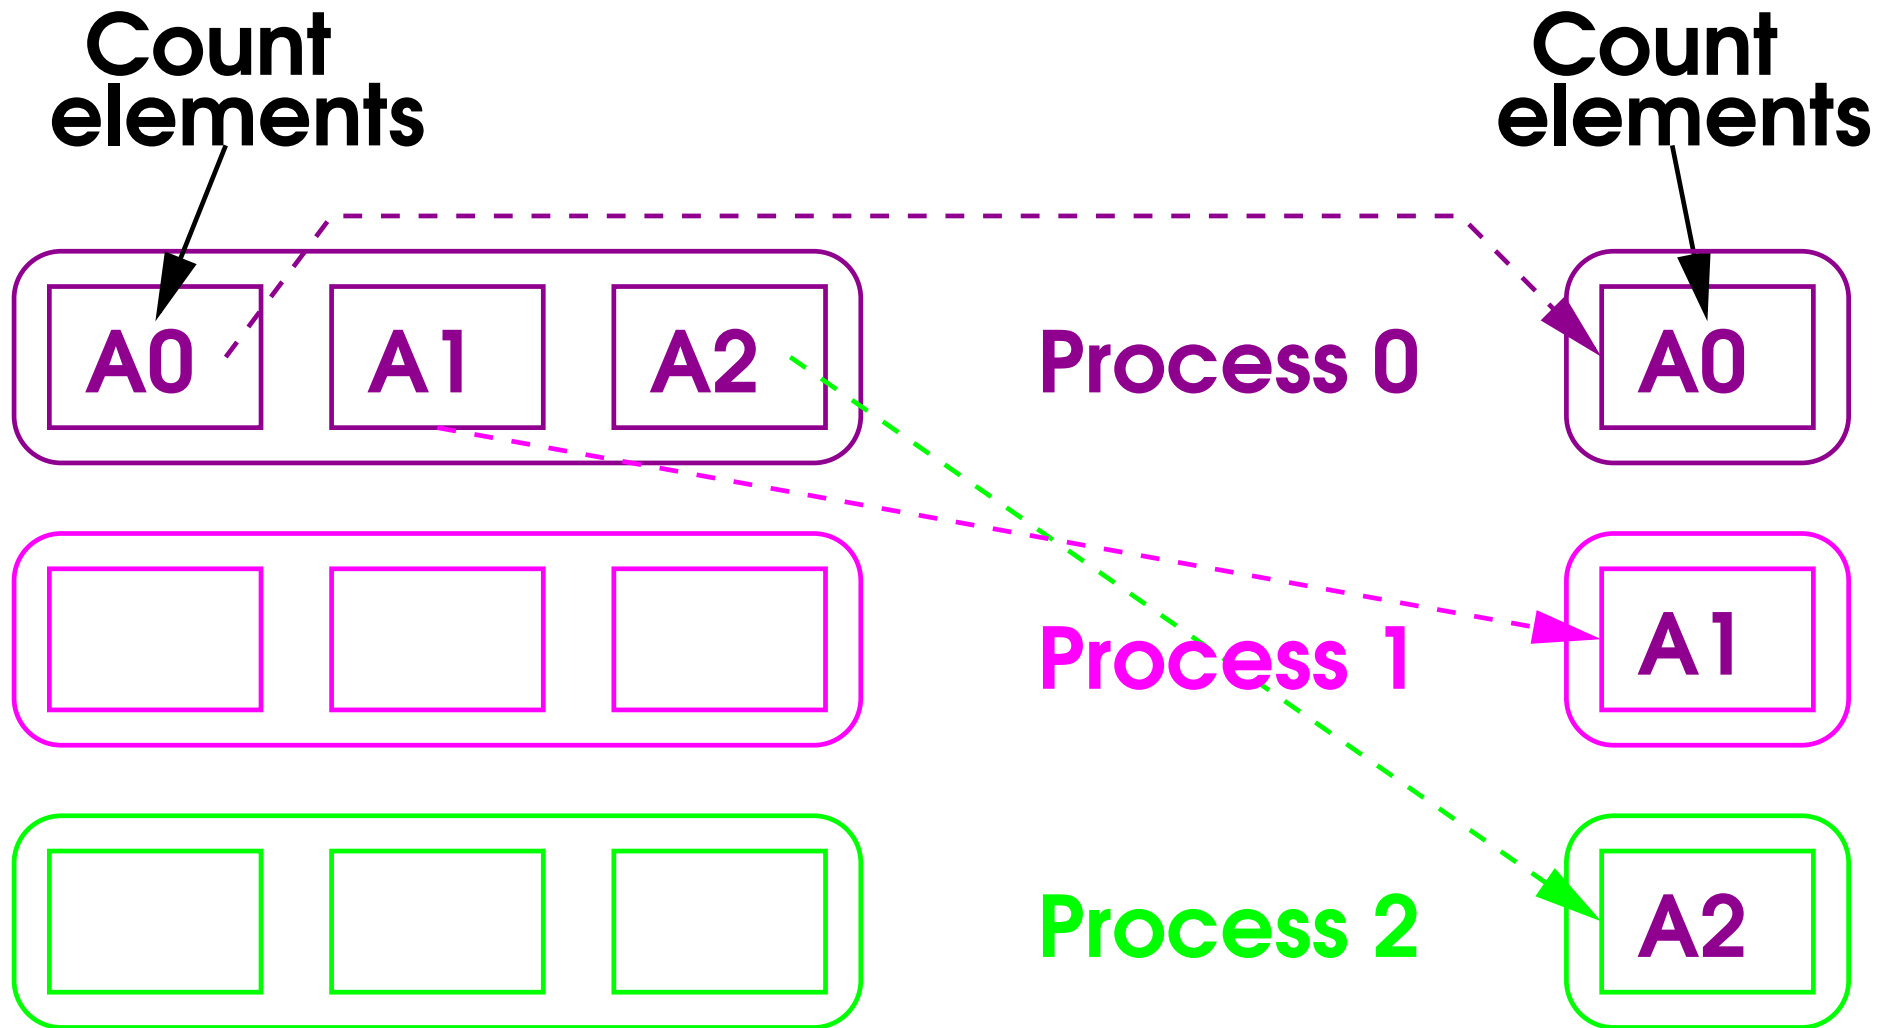
```
vector < double > sendbuf ( 30 * 100 ) , recvbuf ( 100 ) ;  
int root = 3 ;  
error = MPI_Scatter (   
    & sendbuf . front ( ) , 100 , MPI_DOUBLE ,  
    & recvbuf . front ( ) , 100 , MPI_DOUBLE ,  
    root , MPI_COMM_WORLD )
```

Remember that only the **contents** are **contiguous**

Do **NOT** create multiple buffers like this:

```
array< array< double , 100 > , 30 > sendbuf ;
```


Scatter



Hiatus

That is the **basic principles** of collectives

Now might be a good time to do some examples
The first few questions cover the material so far

After that, we cover **datatypes** more thoroughly
And describe more of the **collectives**

Fortran Datatypes (1)

Recommended **datatypes**:

MPI_CHARACTER (\equiv **CHARACTER**(LEN=1))

MPI_LOGICAL

MPI_INTEGER

MPI_REAL

MPI_DOUBLE_PRECISION

MPI_COMPLEX

MPI_DOUBLE_COMPLEX

I.e. **COMPLEX**(KIND=KIND(0.0D0))

Fortran Datatypes (2)

Fortran 90 parameterized types are also supported
`REAL(KIND=SELECTED_REAL_KIND(15,300))`

Details of that mechanism is given later

For use from Fortran, that's all I recommend
There are some more built-in datatypes, though

`MPI_PACKED`, for MPI derived datatypes

`MPI_BYTE` (uninterpreted 8-bit bytes)

What you can do with these is a bit restricted

Other Fortran Datatypes

And you should definitely avoid these

MPI_INTEGER1 MPI_REAL2
MPI_INTEGER2 MPI_REAL4
MPI_INTEGER4 MPI_REAL8

MPI_<type>N translates to <type>*N

That notation is non-standard and outmoded

- It doesn't mean the size in bytes!

E.g. REAL*2 works only on Cray vector systems

C Datatypes (1)

MPI_CHAR is for **char**, meaning **characters**

Don't use it for small integers and arithmetic

Recommended **integer datatypes**:

MPI_UNSIGNED_CHAR

MPI_SIGNED_CHAR

MPI_SHORT

MPI_UNSIGNED_SHORT

MPI_INT

MPI_UNSIGNED (**not** **MPI_UNSIGNED_INT**)

MPI_LONG

MPI_UNSIGNED_LONG

C Datatypes (2)

Recommended floating-point datatypes:

`MPI_FLOAT`

`MPI_DOUBLE`

`MPI_LONG_DOUBLE`

For use from `C`, I recommend one more

`MPI_BYTE` (uninterpreted 8-bit bytes)

What you can do with these is a bit restricted

C++ Datatypes

Warning: these all depend on **MPI 3.0**

Most implementations won't have them yet

MPI_CXX_BOOL

MPI_CXX_FLOAT_COMPLEX

MPI_CXX_DOUBLE_COMPLEX

MPI_CXX_LONG_DOUBLE_COMPLEX

They all correspond to the obvious **C++** type

C99 and Complex

There are types for C99 `_Complex`, if you use it
But I don't advise using that (or most of C99)

C99 `_Complex` may not be compatible with C++
And WG14 have now made `_Complex` optional

`MPI_C_FLOAT_COMPLEX`

`MPI_C_COMPLEX` is a synonym

`MPI_C_DOUBLE_COMPLEX`

`MPI_C_LONG_DOUBLE_COMPLEX`

Other C Datatypes

I don't recommend the other built-in datatypes

`MPI_LONG_LONG_INT` (note the name)

Needs `C99` and optional, anyway

`MPI_UNSIGNED_LONG_LONG`

It need `C99` and is optional, anyway

`MPI_WCHAR` (whatever `C wchar_t` is)

No useful specification in `C90`, `C99` or `C++`

`MPI_PACKED`, for MPI derived datatypes

There is no support for `C99`'s new integer types

- Ask me offline why that is a **Good Thing**

C++, C99 and Complex

Some types for C++, including `bool` and `complex`

And quite a lot of types for C99 extensions

Omitted, for simplicity, because there are problems
The handouts describe the situation in more detail

Gather

Gather is precisely the converse of scatter

- Just change `MPI_Scatter` to `MPI_Gather`

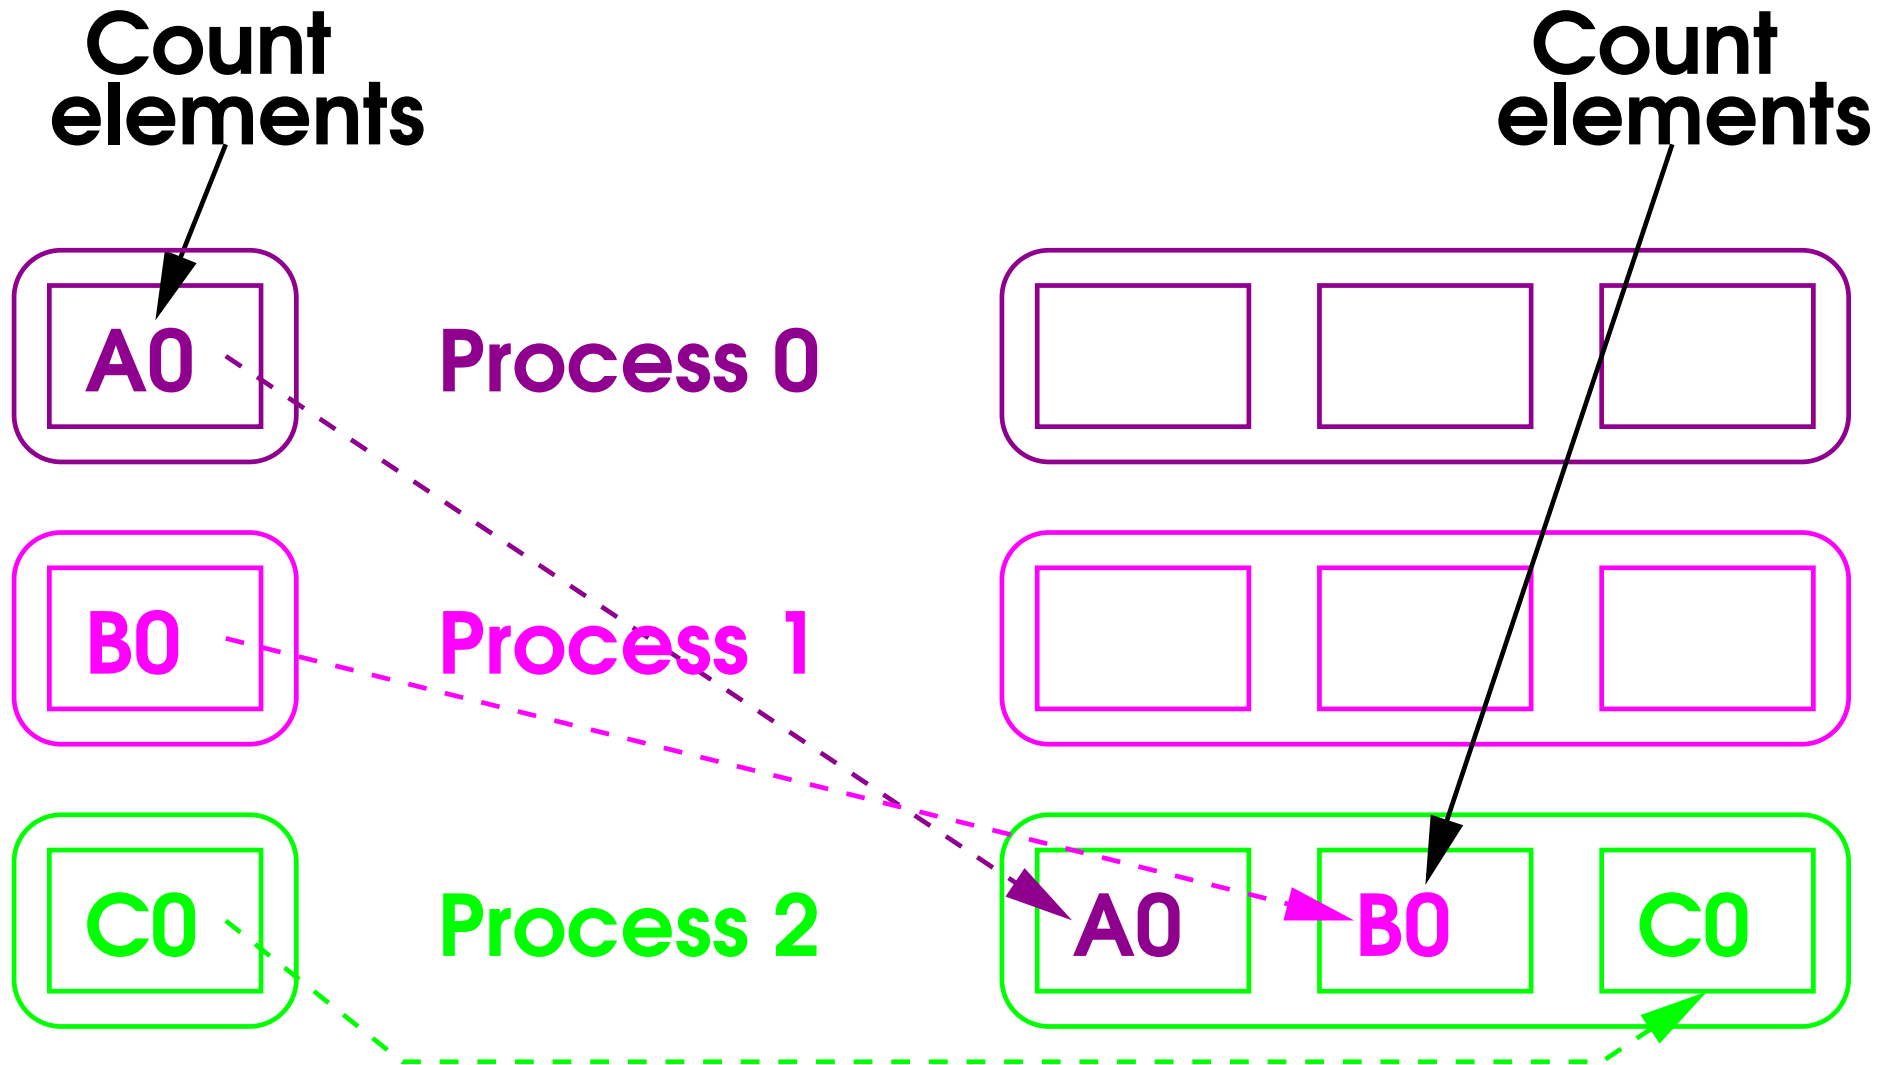
Of course, the array sizes need changing

- It is the receive buffer that needs to be bigger

The send buffer is used on all processes

The receive buffer is used only on the root

Gather

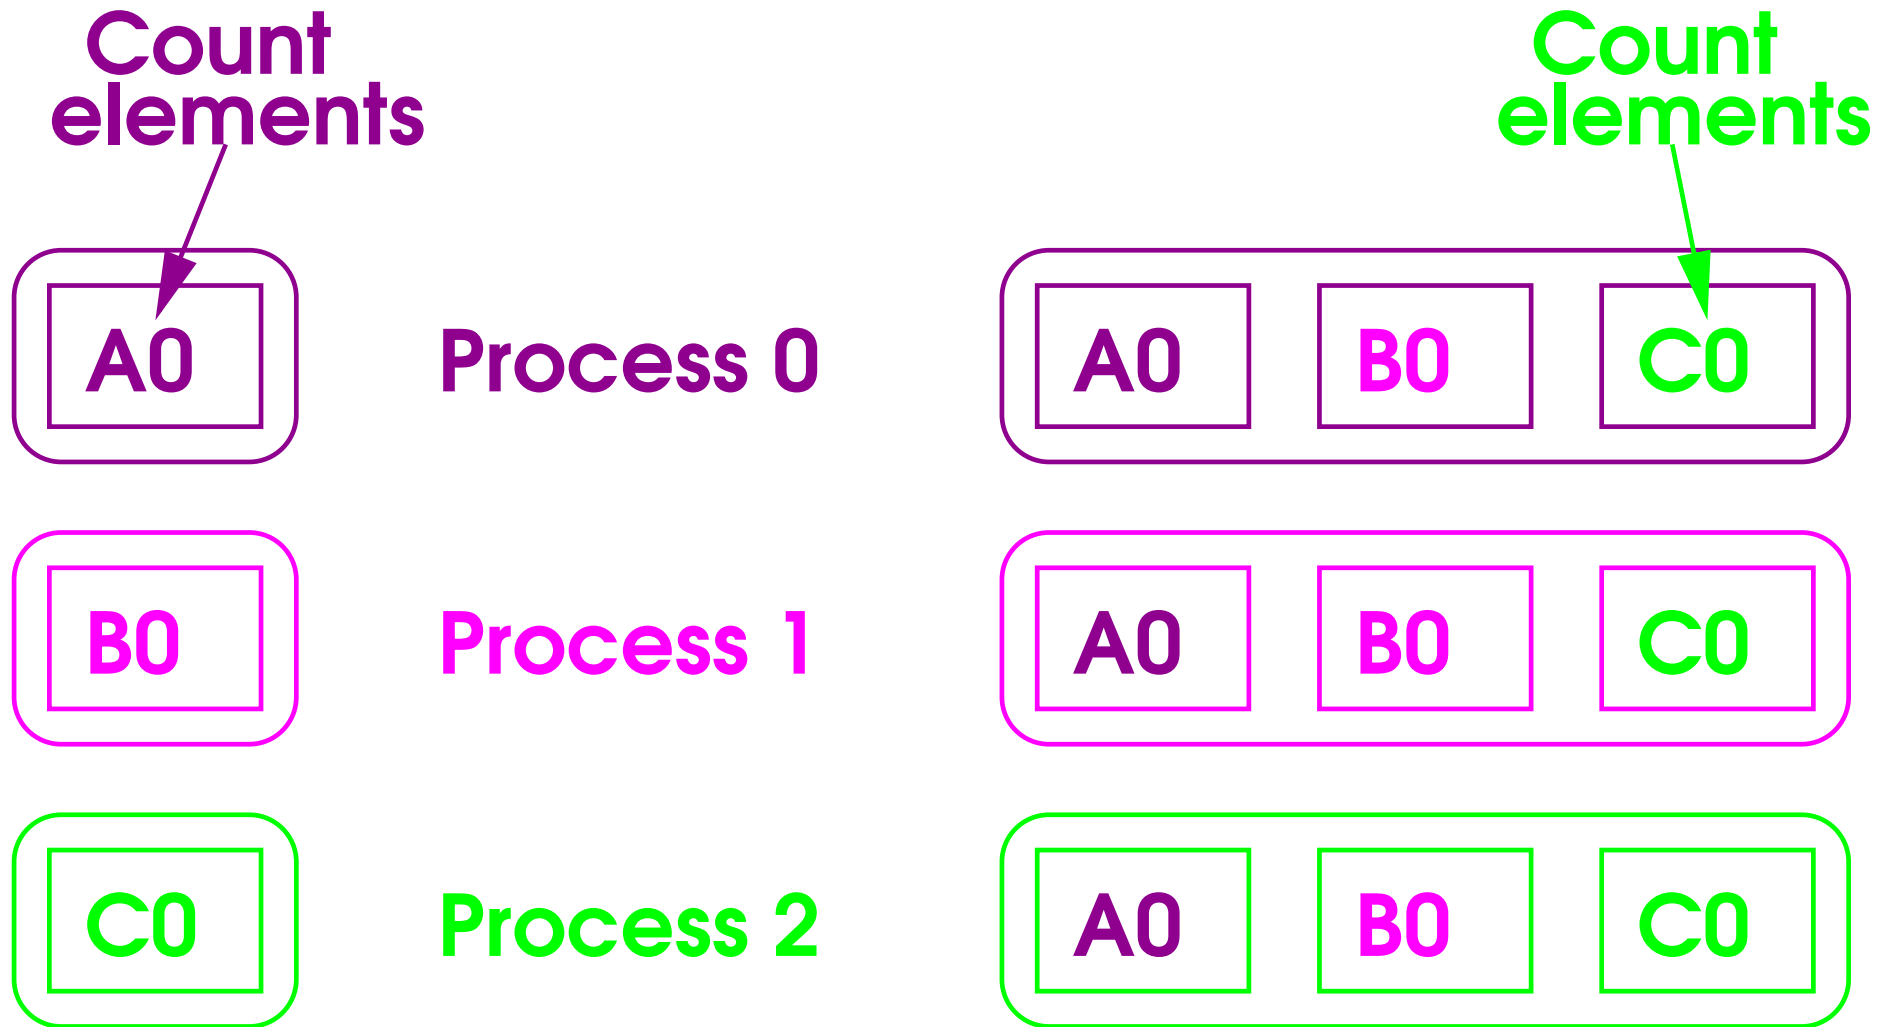


Allgather (1)

You can **gather** data and then **broadcast** it
The interface is very similar, with one difference

- This is now a **symmetric** operation
So has no **argument** specifying the **root** process
- Change **MPI_Gather** to **MPI_Allgather**
And remove the **root** process **argument**, of course
- The **receive buffer** is now used on all **processes**

Allgather



Allgather (2)

Fortran example:

```
REAL(KIND=KIND(0.0D0)) ::      &  
    sendbuf ( 100 ) , recvbuf ( 100 , 30 )  
INTEGER :: error  
CALL MPI_Allgather (      &  
    sendbuf , 100 , MPI_DOUBLE_PRECISION ,      &  
    recvbuf , 100 , MPI_DOUBLE_PRECISION ,      &  
    MPI_COMM_WORLD , error )
```


Allgather (3)

C example:

```
double sendbuf [ 100 ] , recvbuf [ 30 ] [ 100 ] ;  
int error ;  
error = MPI_Allgather (   
    sendbuf , 100 , MPI_DOUBLE ,  
    recvbuf , 100 , MPI_DOUBLE ,  
    MPI_COMM_WORLD )
```

Alltoall

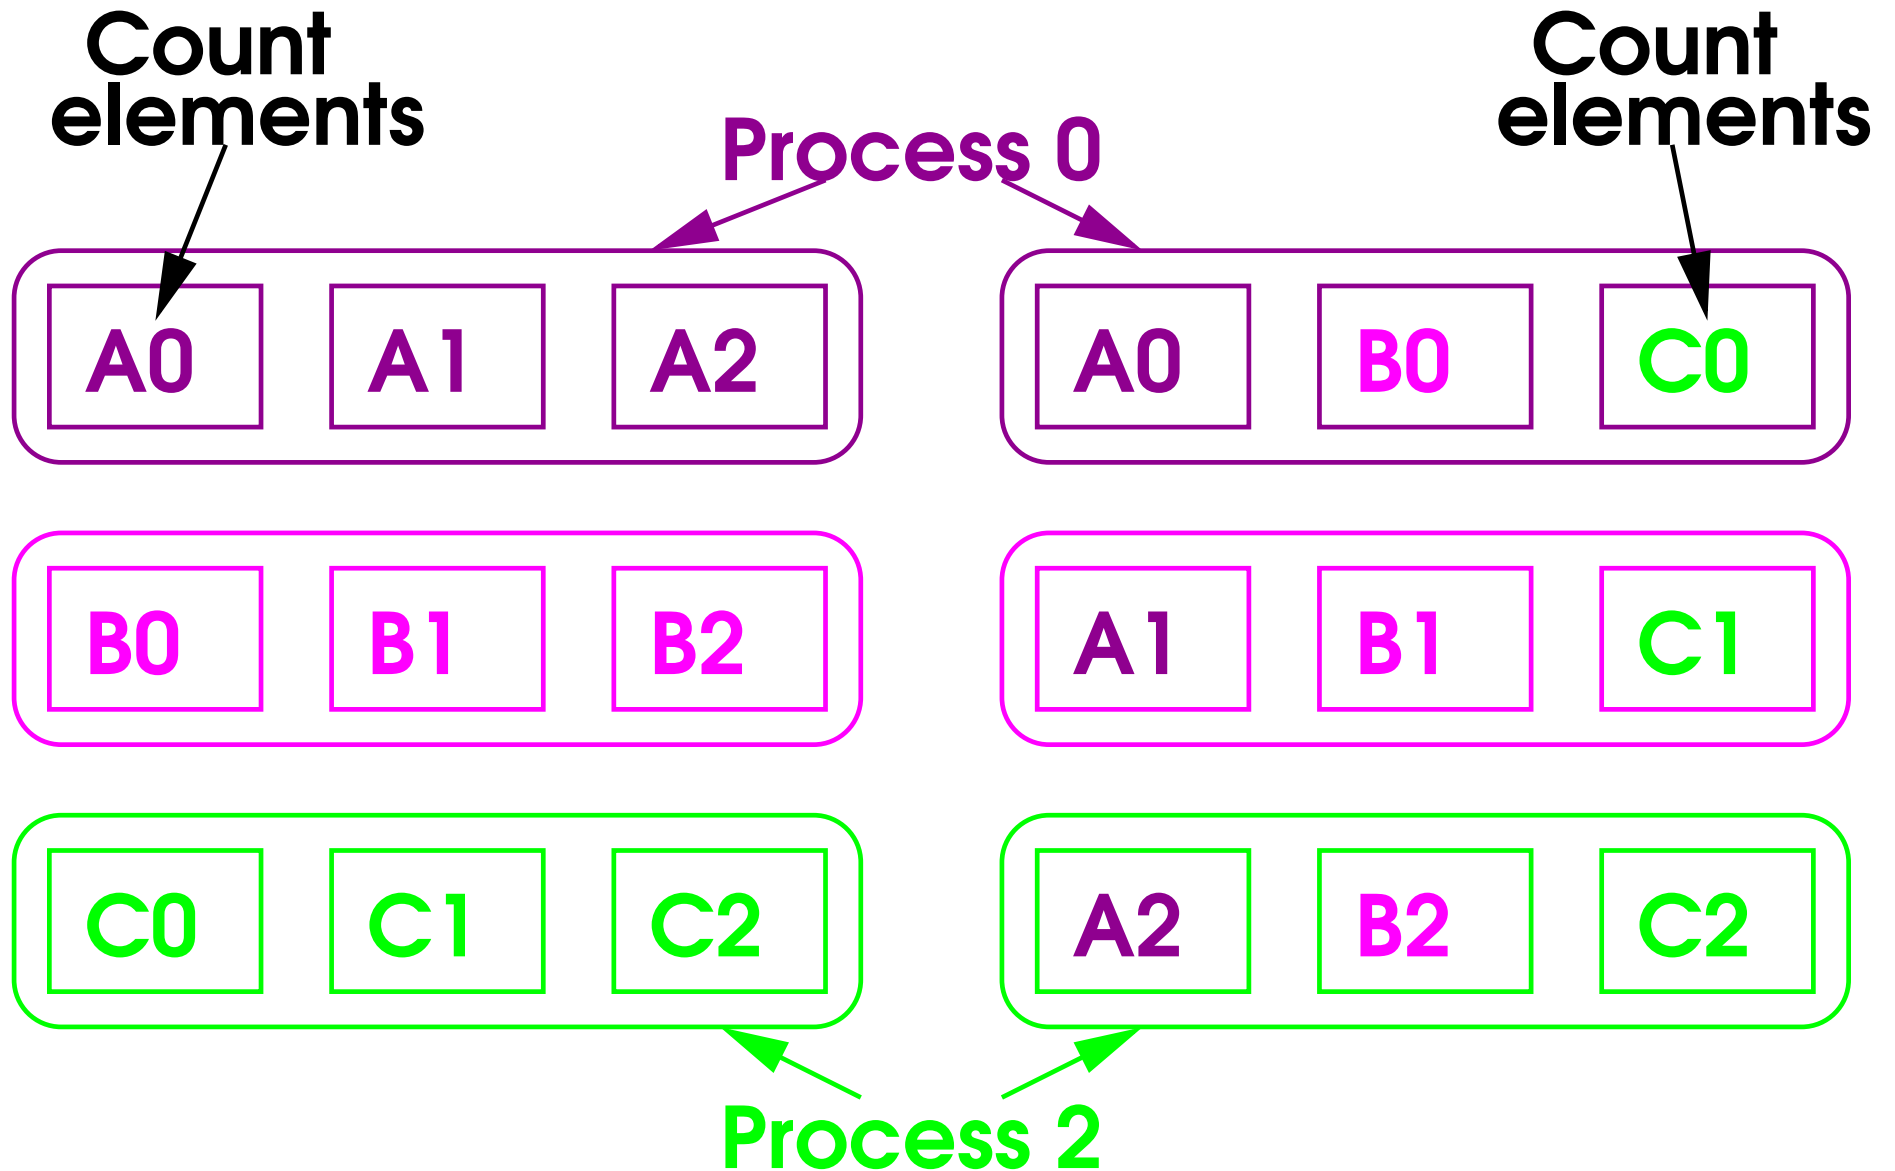
You can do a composite **gather/scatter** operation
Essentially the same interface as **MPI_Allgather**

- Just change **MPI_Allgather** to **MPI_Alltoall**
- Now, **both buffers** need to be bigger

Think of this as a sort of **parallel transpose**
Used when implementing **matrix transpose**

- It's very powerful – a key for **performance**

Alltoall



Global Reductions (1)

One of the basic **parallelisation** primitives

Start with a normal **gather** operation

Then **sum** the values over all **processes**

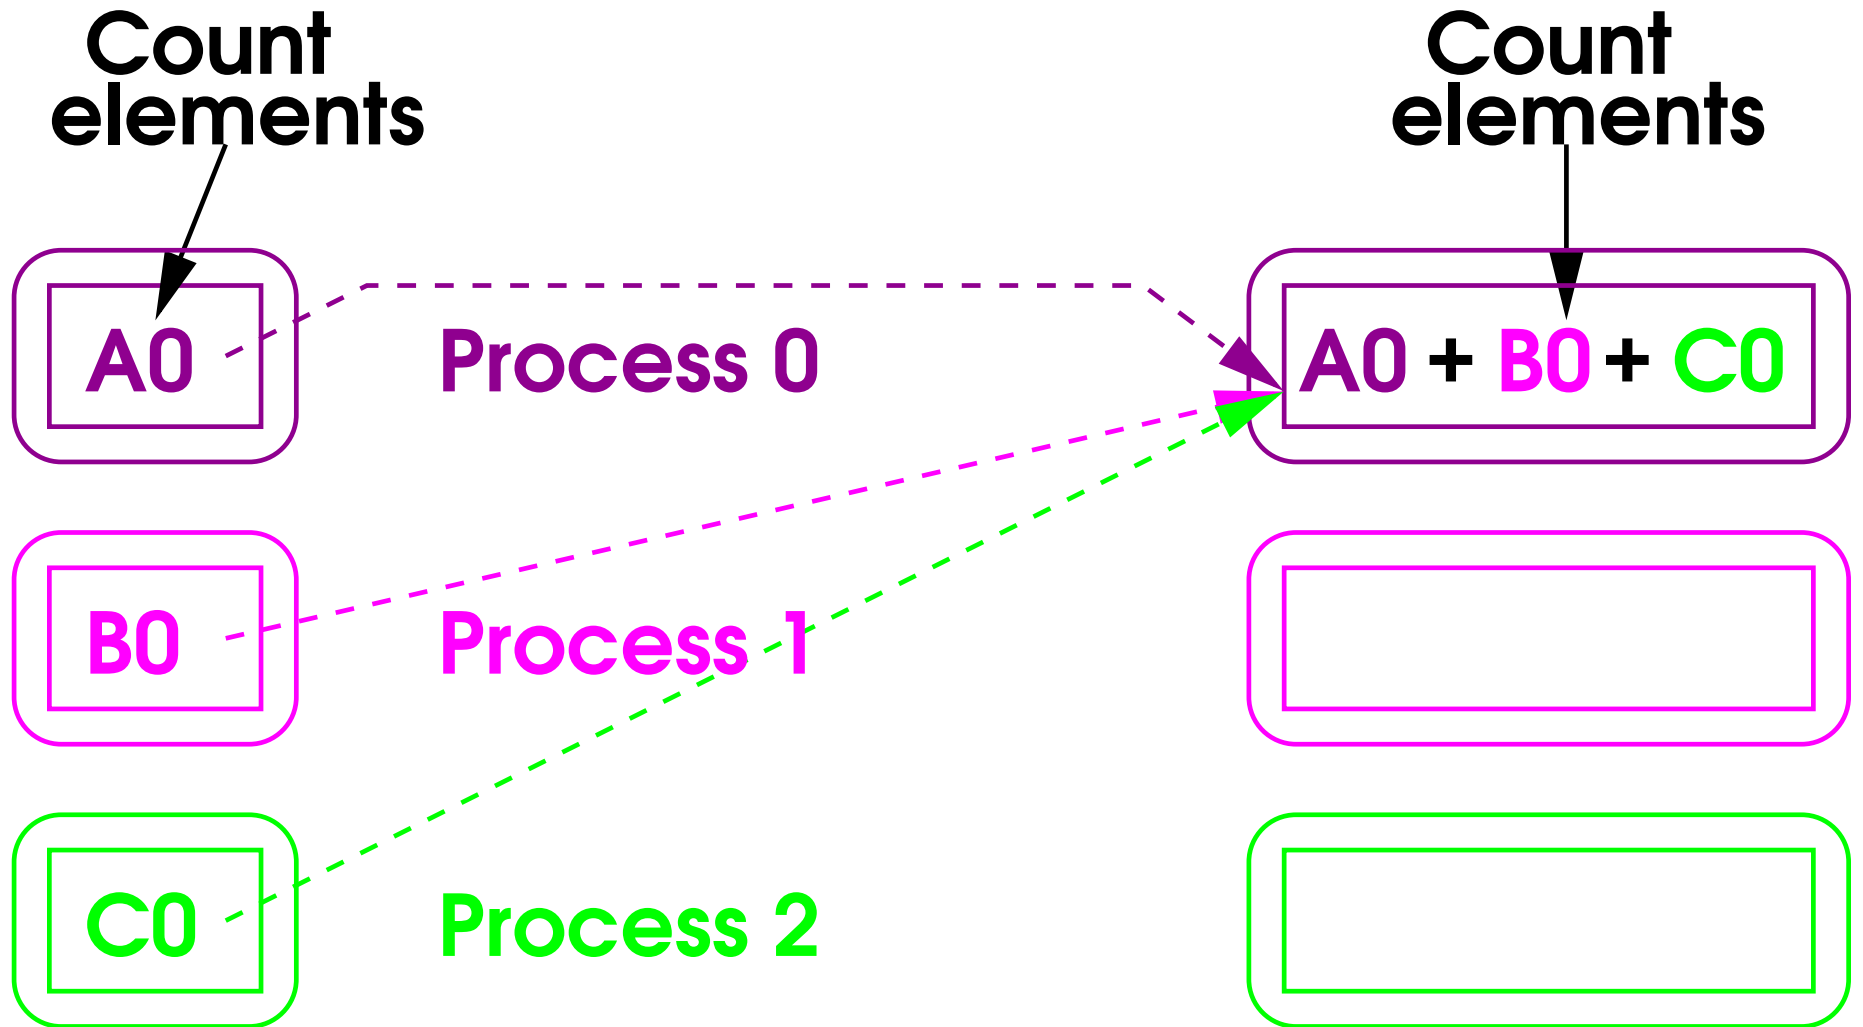
Often can be implemented much more efficiently

- **Summation** is not the only **reduction**

Anything that makes mathematical sense

All of the standard ones are provided

Reduce



Global Reductions (2)

- It specifies the **datatype** and **count** once
Not separately for the **source** and **result**
It makes no sense to do that, so MPI doesn't

- Does **not reduce** over the **vector**
The **count** is the size of the **result**, too
It sums the values for each **index** separately

You have to **reduce** over the **vector** yourself

- Doing it **beforehand** is more efficient

Reduce Result

Process 0

A0

B0

C0

Process 1

A1

B1

C1

Process 2

A2

B2

C2

Result

A0+A1+A2

B0+B1+B2

C0+C1+C2

Reduce (2)

Fortran example:

```
REAL(KIND=KIND(0.0D0)) ::      &  
    sendbuf ( 100 ) , recvbuf ( 100 )  
INTEGER , PARAMETER :: root = 3  
INTEGER :: error  
CALL MPI_Reduce ( sendbuf , recvbuf ,      &  
    100 , MPI_DOUBLE_PRECISION ,      &  
    MPI_SUM , root , MPI_COMM_WORLD , error )
```


Reduce (3)

C example:

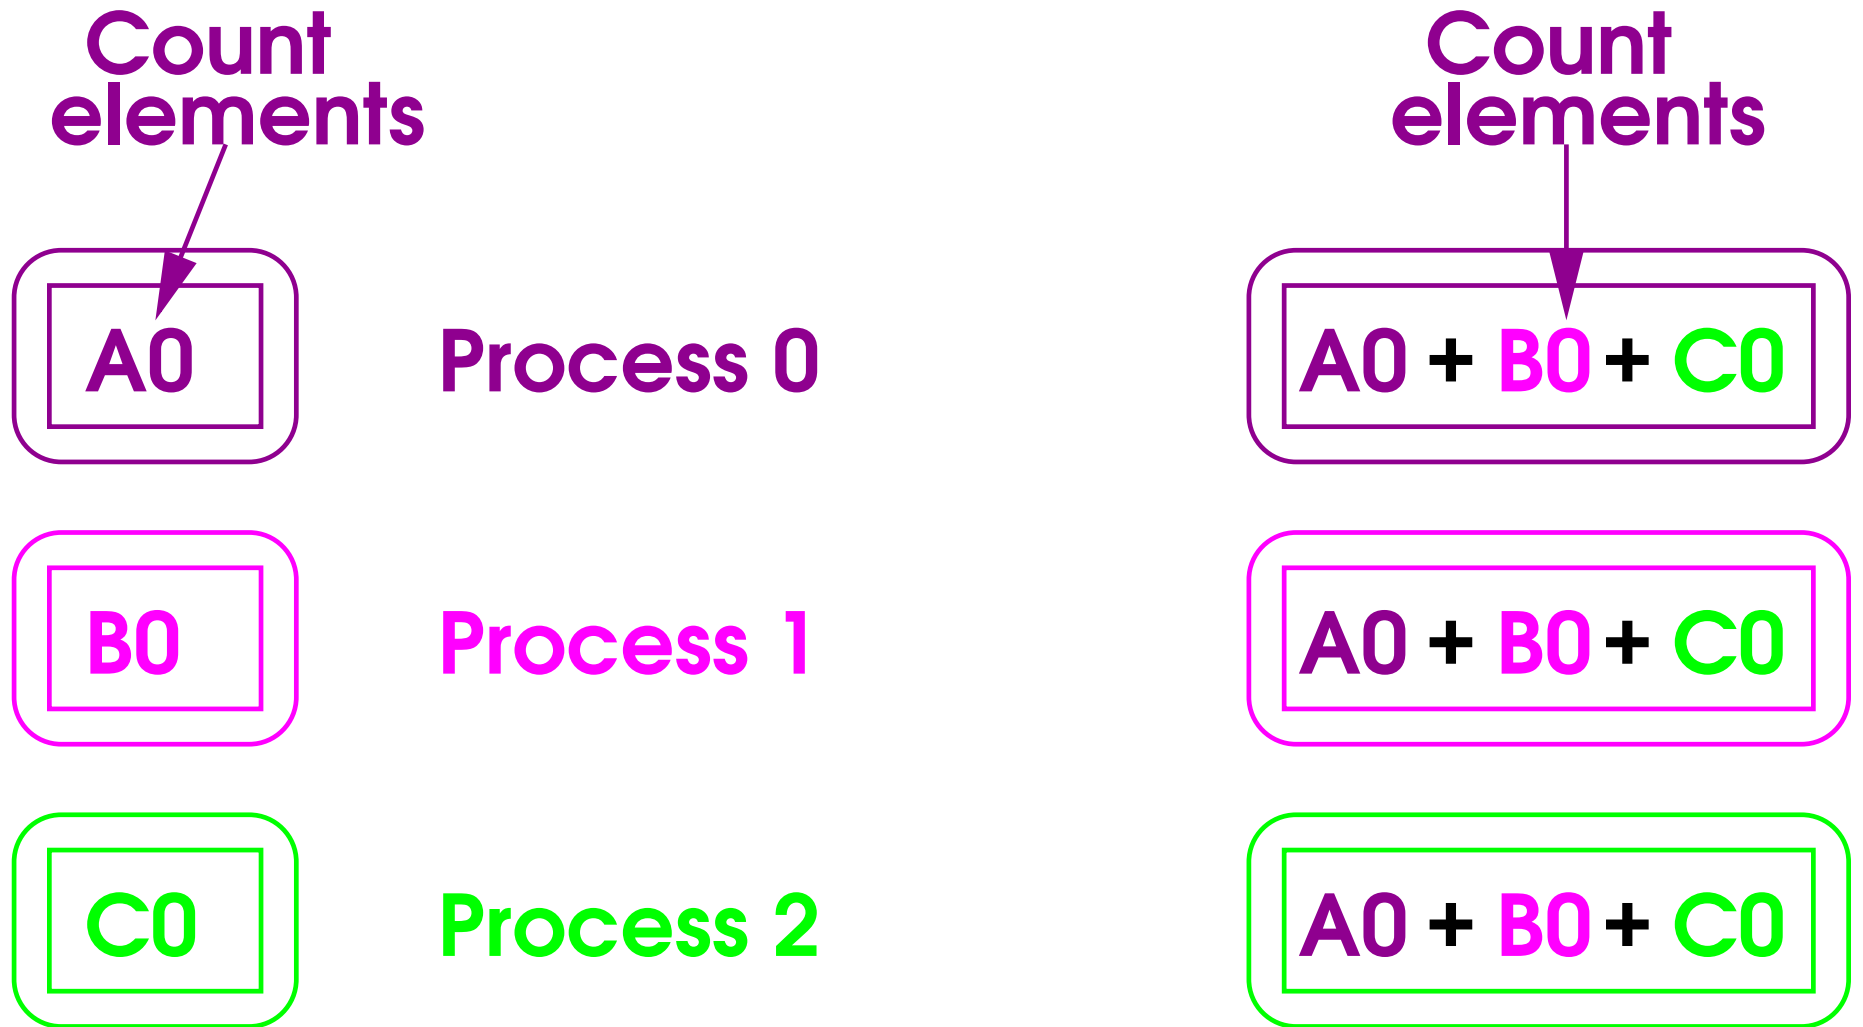
```
double sendbuf [ 100 ] , recvbuf [ 100 ] ;  
int root = 3 , error ;  
error = MPI_Reduce ( sendbuf , recvbuf ,  
                    100 , MPI_DOUBLE , MPI_SUM , root ,  
                    MPI_COMM_WORLD )
```

Allreduce

You can **reduce** data and then **broadcast** it
Again, the interface is essentially identical

- This is now a **symmetric** operation
So has no **argument** specifying the **root** process
- Just change **MPI_Reduce** to **MPI_Allreduce**
And remove the **root** process **argument**, of course
- The **receive buffer** is now used on all **processes**

Allreduce



Reduction Operations (1)

Same rules for all precisions of **number**

MPI_MIN **integer** or **real** minimum

MPI_MAX **integer** or **real** maximum

MPI_SUM **integer**, **real** or **complex** sum

MPI_PROD **integer**, **real** or **complex** product

Note there are no reductions on **character** data

Reduction Operations (2)

Boolean is **int** in **C** and **LOGICAL** in **Fortran**

The supported values are **only** **True** and **False**

You can also perform **bitwise** operations on **integers**

MPI LAND	Boolean AND
MPI LOR	Boolean OR
MPI LXOR	Boolean Exclusive OR
MPI BAND	integer bitwise AND
MPI BOR	integer bitwise OR
MPI BXOR	integer bitwise Exclusive OR

More on Collectives

There is a little more to say on **collectives**
But that's quite enough for now

The above has covered all of the essentials
The remaining aspects to cover are:

- **Fortran** parameterised types
- Searching as a **reduction**
- Using more flexible buffer layouts
- Using **collectives** efficiently

Practicals

There are a lot of exercises on the above
Will take you through almost all aspects

- Each one should need very little editing/typing
You can start from a previous one as a basis

PLEASE check you understand the point
And that you get the same answers as are provided
And that you understand what it is doing and why

- They are pointless if you do them mechanically