

Programming with MPI

Point-to-Point Transfers

Nick Maclaren

nmm1@cam.ac.uk

May 2008

Digression

Most books and courses teach **point-to-point** first
And then follow up by teaching **collectives**

This course hasn't — why not?

- **Point-to-point** is hard to use **correctly**
I usually make a complete mess of it, first time
See **Hoare's Communicating Sequential Processes**
Hoare designed **BSP** based on his experience!

After all, who programs in **assembler** nowadays?
Point-to-point is the **assembler-level** interface

Using Point-to-Point

- Above all, **KISS** – Keep It Simple and Stupid
- Design proper **primitives**, don't just code

Simplest to use one of two **design models** for that:

- Your own **collective** – see later
 - **Two** processes, doing nothing else
-
- It's easiest if your **primitives** don't overlap
- Can separate by **barriers** and **debug** separately
- Almost essential for **tuning** – see later

Envelopes

Think of **point-to-point** as a sort of **Email**
Like that, **messages** come in **envelopes**

MPI's **envelopes** contain the following:

- The **source** process
- The **destination** process
- The **communicator**
- An identifying **tag**

One of the **first two** is the **calling process**
The others are specified in the **arguments**

Receive Status (1)

A **receive** action returns a **status**

This contains the following:

- The **source** process
- The identifying **tag**
- Other, **hidden**, information

Already know the **communicator** and **destination**

A **function** to extract the **message size**

Receive Status (2)

In **C**, the **status** is a typedef structure **MPI_Status**

In **Fortran**, it is an integer array
INTEGER, DIMENSION(MPI_STATUS_SIZE)

- You **declare** these yourself, as normal
Including in **static** memory or on the **stack**
- They are **not** like **communicators**
You **don't** call MPI to **allocate** and **free** them

Receive Status (3)

For now, you can largely ignore the status
You don't need to look at it for very simple use

- In **MPI 1**, had to provide the **argument**
This is the form that I shall use in examples
- **MPI 2** allowed you to not provide it
I don't recommend doing that, in general

The Simplest Use

Assume `communicator` is `MPI_COMM_WORLD`

The `tag` is needed only for advanced use
Quite useful for added `checking`, though

So it's only the `destination` and `source`
The `latter` is set automatically for `send`!
And the `former` is for `receive`!

The functions are `MPI_Send` and `MPI_Recv`

Fortran Example (1)

```
REAL(KIND=KIND(0.0D0)) :: buffer ( 100 )
INTEGER :: myrank , error
INTEGER , PARAMETER :: from = 2 , to = 3 ,      &
    tag = 123

CALL MPI_Rank ( myrank , error )
IF ( myrank == from ) THEN
    CALL MPI_Send ( buffer , 100 ,      &
        MPI_DOUBLE_PRECISION , to , tag ,      &
        MPI_COMM_WORLD , error )
END IF
```

Fortran Example (2)

```
REAL(KIND=KIND(0.0D0)) :: buffer ( 100 )
INTEGER :: myrank , error , status ( MPI_STATUS_SIZE )
INTEGER , PARAMETER :: from = 2 , to = 3 ,      &
    tag = 123

CALL MPI_Rank ( myrank , error )
IF ( myrank == to ) THEN
    CALL MPI_Recv ( buffer , 100 ,      &
        MPI_DOUBLE_PRECISION , from , tag , &
        MPI_COMM_WORLD , status , error )
END IF
```

C Example (1)

```
double buffer [ 100 ] ;  
int myrank , from = 2 , to = 3 , tag = 123 , error ;  
  
error = MPI_Rank ( & myrank ) ;  
if ( myrank == from )  
    error = MPI_Send ( buffer , 100 , MPI_DOUBLE ,  
                        to , tag , MPI_COMM_WORLD ) ;
```

C Example (2)

```
double buffer [ 100 ] ;  
MPI_Status status ;  
int myrank , from = 2 , to = 3 , tag = 123 , error ;  
  
error = MPI_Rank ( & myrank ) ;  
if ( myrank == to )  
    error = MPI_Recv ( buffer , 100 , MPI_DOUBLE ,  
        from , tag , MPI_COMM_WORLD , & status ) ;
```

Beyond That

Trivial as that is, it's enough to cause trouble
There are some examples on how that can happen

And it's not enough for all real programs
MPI provides lots of knobs, bells and whistles

- You should use only what you need
Don't use something because it looks cool
- You need to know what can be done
When you need something extra, look it up

Blocking (1)

Receive will block until a matching **send**

If one is never posted, it will hang indefinitely

Send **may** block until a matching **receive**

Or it **may** copy the message and return

and MPI will transfer it in due course

Unspecified and up to the **implementation**

May vary between messages, or phase of the moon

- Correct MPI programs will work either way
- You can control that yourself – see later

Blocking (2)

Processes **A** and **B** want to swap data

Both **send** the existing value, and then **receive**?
It will **sometimes** work and sometimes hang

Process **A**

send to **B**

Process **B**

send to **A**

Both **may** wait until transfers **received**

receive from **B** **receive** from **A**

Blocking (3)

In that case, it's trivial to avoid

- If $A < B$, A sends first and receives second
And B receives first and sends second

And conversely if $A > B$

Complicated transfer graphs are easy to get wrong
MPI provides several ways to avoid the problem
Use whichever is simplest for your purposes

Transfer Modes (1)

`MPI_Ssend` is **synchronous** (**will** block)
returns when the message has been **received**

`MPI_Bsend` is **buffered** (**won't** block)
so the swap example above will never hang

- Exactly the same usage as for `MPI_Send`
`MPI_Send` simply calls one or the other

Generally, don't use either of them
Both have important, but advanced, uses

Transfer Modes (2)

A **synchronous send** avoids a separate handshake
Completing the call acknowledges **receipt**

- Use it if it avoids an explicit acknowledgement

Buffering is more tricky, surprisingly enough
Sends are **erroneous** if the buffer becomes full

- And the **default buffer size** is **zero**!

But **exceeding** it is **undefined behaviour**!

Using buffering is covered later

Composite Send and Receive (1)

- There is a composite **send** and **receive**
Will do the in the right order to avoid **deadlock**
Can also match ordinary **send** and **receive**
 - It also has a form that updates in place
Sends buffer and then **receives** into it
That may involve extra copying, of course
- Use these if they are what you want to do
They aren't likely to be any more efficient

Composite Send and Receive (2)

Fortran example:

```
REAL(KIND=KIND(0.0D0)) ::      &  
    putbuf ( 100 ) , getbuf ( 100 )  
INTEGER :: error , status ( MPI_STATUS_SIZE )  
INTEGER , PARAMETER :: from = 2 , to = 3 ,      &  
    fromtag = 123 , totag = 456  
  
CALL MPI_Sendrecv ( putbuf , 100 ,      &  
    MPI_DOUBLE_PRECISION , to , totag ,      &  
    getbuf , 100 , MPI_DOUBLE_PRECISION ,      &  
    from , fromtag ,      &  
    MPI_COMM_WORLD , status , error )
```

Composite Send and Receive (3)

Fortran in place example:

```
REAL(KIND=KIND(0.0D0)) :: buffer ( 100 )  
INTEGER :: error , status ( MPI_STATUS_SIZE )  
INTEGER , PARAMETER :: from = 2 , to = 3 ,      &  
    fromtag = 123 , totag = 456  
  
CALL MPI_Sendrecv_replace (      &  
    buffer , 100 , MPI_DOUBLE_PRECISION ,      &  
    to , totag , from , fromtag ,      &  
    MPI_COMM_WORLD , status , error )
```

C Example

```
double putbuf [ 100 ] , getbuf [ 100 ] , buffer [ 100 ] ;  
MPI_Status status ;
```

```
int from = 2 , to = 3 , fromtag = 123 , totag = 456 ,  
    error ;
```

```
error = MPI_Sendrecv (  
    putbuf , 100 , MPI_DOUBLE , to , totag ,  
    getbuf , 100 , MPI_DOUBLE , from , fromtag ,  
    MPI_COMM_WORLD , & status ) ;
```

```
error = MPI_Sendrecv_replace (  
    buffer , 100 , MPI_DOUBLE , to , totag ,  
    from , fromtag , MPI_COMM_WORLD , & status ) ;
```

Unknown Message Size (1)

The **send** and **receive** sizes need not match

- It is an error if the **receive** is smaller

Only the **send count** values are updated

E.g. **sending** 30 items and **receiving** 100 items
will leave the last 70 items unchanged

- But there is a better way to do this

Allows **receiving** truly unknown size messages

This is where you start to use the **status**

Unknown Message Size (2)

- Can **accept** the message with **MPI_Probe**
Calling it **probe** is a bit of a misnomer
It **accepts** the message and updates the **status**
But it doesn't transfer the data anywhere
- You discover the **size** with **MPI_Get_count**
Then you can allocate a suitable buffer
MPI_Get_count needs the **datatype**
Allows for **conversion**, not covered here
- Lastly, you **receive** the message as normal

Fortran Example

```
REAL(KIND=KIND(0.0D0)) ,      &  
    ALLOCATABLE :: buffer ( : )  
INTEGER :: error , count ,      &  
    status ( MPI_STATUS_SIZE )  
INTEGER , PARAMETER :: from = 2 , tag = 123  
  
CALL MPI_Probe ( from , tag ,      &  
    MPI_COMM_WORLD , status , error )  
CALL MPI_Get_count ( status ,      &  
    MPI_DOUBLE_PRECISION , count , error )  
ALLOCATE ( buffer ( count ) )  
CALL MPI_Recv ( buffer , count ,      &  
    MPI_DOUBLE_PRECISION , ...
```

C Example

```
double * buffer ;  
int from = 2 , tag = 123 , error , count ;  
MPI_Status status ;  
  
error = MPI_Probe ( from , tag ,  
    MPI_COMM_WORLD , & status ) ;  
error = MPI_Get_count ( & status ,  
    MPI_DOUBLE , & count ) ;  
buffer = malloc ( sizeof ( double ) * count ) ;  
if ( buffer == NULL ) . . . ;  
error = MPI_Recv ( buffer , count , MPI_DOUBLE ,  
    from , tag , MPI_COMM_WORLD , & status ) ;
```

Checking for Messages (1)

- Real **probe** function is called **MPI_Iprobe**
It returns immediately even if no matching message

An extra **Boolean** argument saying if there is one

- If there is one, it behaves just like **MPI_Probe**
- If there isn't one, the **status** is not updated

It's so similar, shall show only the actual differences

Checking for Messages (2)

Fortran example:

```
LOGICAL :: flag
```

```
CALL MPI_Iprobe ( from , tag ,      &  
                MPI_COMM_WORLD , flag , status , error )
```

C example:

```
int flag ;
```

```
error = MPI_Iprobe ( from , tag ,  
                  MPI_COMM_WORLD , & flag , & status ) ;
```

Wild Cards (1)

- You can **accept** messages from any **process**
Just use **MPI_ANY_SOURCE** for **from**

The actual **source** is stored in the **status**
using the name **MPI_SOURCE**

Fortran example: **status(MPI_SOURCE)**

C example: **status . MPI_SOURCE**

- Be warned – your footgun is now loaded

Wild Cards (2)

- You can **accept** messages with any **tag**
Just use **MPI_ANY_TAG** for **tag**
Use the name **MPI_TAG** like **MPI_SOURCE**

I advise using the **tag** for **cross-checking**

- It could be a message **sequence number**
 - Or identify the **object** being transferred
 - Or whatever else would help debugging
- On **receipt**, check it is what you **expect**
If it isn't, you can write your own diagnostics
Including as much program state as you want

Fortran Example

```
INTEGER :: error , count , from , tag ,      &  
        status ( MPI_STATUS_SIZE )  
  
CALL MPI_Probe ( MPI_ANY_SOURCE ,      &  
                MPI_ANY_TAG , MPI_COMM_WORLD ,      &  
                status , error )  
CALL MPI_Get_count ( status ,      &  
                   MPI_DOUBLE_PRECISION , count , error )  
from = status ( MPI_SOURCE )  
tag = status ( MPI_TAG )
```

C Example

```
int error , from , tag , count ;  
MPI_Status status ;  
  
error = MPI_Probe ( MPI_ANY_SOURCE ,  
                   MPI_ANY_TAG , MPI_COMM_WORLD ,  
                   & status ) ;  
error = MPI_Get_count ( & status ,  
                       MPI_DOUBLE , & count ) ;  
from = status . MPI_SOURCE ;  
tag = status . MPI_TAG ;
```


Message Ordering (1)

Each **process** has a **FIFO** receipt (**queue**)
Incoming messages never overtake each other

Every **probe** and **receive** match in queue order
First message that satisfies **all** of the constraints

Probe and **receive** get same message if

- There has been no intervening **receive**
- Same **communicator**, **source** and **tag**

Other safe usages, too, but that one is easy

Message Ordering (2)

If you **probe** using **wild cards**, you can also extract the **source** and **tag** from **status** and then use **those values** in the **receive**

If process **A** does multiple **sends** to process **B** those messages **arrive** in the **same order**

- **No ordering** if **sender** or **receiver** differ
- And messages can be **delayed** considerably

Tag Warning

The main purpose of **tags** is not for **checking**
It's to allow **independent** communication paths

Many books and Web pages will describe that use
Some will even encourage it

Don't do it

It's the equivalent of cocking your footgun
Using **tags** like that is very hard for **experts**

- I will contradict myself later, under I/O

Buffered Sends (1)

These are trivial to use, but need extra mechanism

- Default buffer size is **implementation dependent** and doesn't even have to be documented!

IBM chose to use **8 bytes** for **poe**

- So you **have** to allocate a buffer first
It's just a block of memory – any type will do

That's really the only extra complexity
And you can usually just make it very big

Buffered Sends (2)

You attach a **single** buffer to a **process**
not a **communicator** — why not?

When you have **finished** doing transfers, **detach** it

- It is used for **scratch space** by MPI in between

Best to set **immediately** after **MPI_Init**
And detach **immediately** before **MPI_Finalize**

The MPI standard is (**unusually**) not very clear
Does the **detach** read its **arguments** or not?
I recommend **setting them** before the call anyway

Buffered Sends (3)

When a **buffer** is **in use** by MPI

- Do **NOT** fiddle with it in **ANY** way!
Its use and contents are completely undefined
- Watch out in **garbage-collected** languages
Make sure that the buffer will not **move around**
- Even in **Fortran** and **C**
Make sure that it does not go **out of scope**
Or falls foul of **Fortran copy-in/copy-out**

Allocating a Buffer (1)

Fortran example:

```
INTEGER , PARAMETER :: buffsize = 10000  
CHARACTER :: buffer ( buffsize )  
INTEGER :: oldsize , error
```

```
CALL MPI_Buffer_attach ( buffer , buffsize , error )
```

```
oldsize = buffsize
```

```
CALL MPI_Buffer_detach ( buffer , oldsize , error )
```

Detach returns the values previously stored
I have no idea what this means for **buffer**!

Allocating a Buffer (2)

C example:

```
#define bufsize 10000
void * buffer = malloc ( bufsize ) , * oldbuff;
int oldsize , error ;

error = MPI_Buffer_attach ( buffer , bufsize ) ;

oldbuff = buffer ;
oldsize = bufsize ;
error = MPI_Buffer_detach ( & oldbuff , & oldsize ) ;
```

Note the indirections (&) in **detach**

Detach stores the values previously stored

Use of Buffered Sends (1)

Using them is generally **not** advisable

They usually **hide** problems rather than **fix** them

And they can be quite a lot less efficient

If you have a completely baffling failure
try changing all **sends** to **buffered**

- If that helps, you have a **race condition**
You then must track it down and fix it properly

The other main use is for I/O (see later)

Use of Buffered Sends (2)

You can calculate how much space you need

Constant `MPI_BSEND_OVERHEAD`

Function `MPI_Pack_size`

Function `MPI_Sizeof`

[Only in `Fortran`]

Using those is overkill for almost all programs
This course doesn't describe their use

Epilogue

There is more on **point-to-point** later

Mainly **non-blocking** (**asynchronous**) transfers

But we have covered most of **blocking** transfers

Exercises will try out quite a lot of this

Main one is to code a rotation **collective**

Each **process** sends to its successor

And the last one sends back to the beginning

Practicals

Practicals often use **buffered** or **synchronous** sends
Reason is to **expose** or **hide** cases of **deadlock**

- This is advised **only** when testing
You should normally use **ordinary** sends