

Programming with MPI

More on Datatypes and Collectives

Nick Maclaren

nmm1@cam.ac.uk

May 2008

Less Basic Collective Use

A few important facilities we haven't covered
Less commonly used, but fairly often needed
In particular, one of them can help a lot with I/O

And then how to use **collectives** efficiently
Plus one potentially useful minor feature

Fortran Precisions (1)

Fortran 90 allows selectable precisions

`KIND=SELECTED_INTEGER_KIND(precision)`

`KIND=SELECTED_REAL_KIND(precision[,range])`

Can create a MPI **derived datatype** to match these
Then can use it just like a **built-in datatype**

Surprisingly, it is a **predefined datatype**

Do **NOT** **commit** or **free** it

[Don't worry if that makes no sense to you]

Fortran Precisions (2)

```
INTEGER ( KIND =      &  
          SELECTED_INTEGER_KIND ( 15 ) ) ,      &  
          DIMENSION ( 100 ) :: array  
INTEGER :: root , integertype , error  
  
CALL MPI_Type_create_f90_integer (      &  
          15 , integertype , error )  
CALL MPI_Bcast ( array , 100 ,      &  
          integertype , root ,      &  
          MPI_COMM_WORLD , error )
```

Fortran Precisions (3)

REAL and **COMPLEX** are very similar

```
REAL ( KIND =      &  
      SELECTED_REAL_KIND ( 15 , 300 ) ) ,      &  
      DIMENSION ( 100 ) :: array  
CALL MPI_Type_create_f90_real (      &  
      15 , 300 , realtype , error )
```

```
COMPLEX ( KIND =      &  
      SELECTED_REAL_KIND ( 15 , 300 ) ) ,      &  
      DIMENSION ( 100 ) :: array  
CALL MPI_Type_create_f90_complex (      &  
      15 , 300 , complextype , error )
```

Searching (1)

You can use **global reductions** for **searching**

Bad news: it needs MPI's **derived datatypes**

Good news: there are some useful built-in ones

We do a **reduction** with a composite **datatype**:
(**<value>** , **<index>**)

As with **sum**, we build up from a **binary operator**

Two **built-in** operators:

MPI_MINLOC and **MPI_MAXLOC**

We shall use finding the **minimum** as an example

Searching (2)

$(\langle \text{value_1} \rangle, \langle \text{index_1} \rangle) \& (\langle \text{value_2} \rangle, \langle \text{index_2} \rangle)$

We shall produce a result $(\langle \text{value_x} \rangle, \langle \text{index_x} \rangle)$

If $\langle \text{value_1} \rangle \leq \langle \text{value_2} \rangle$ then

$\langle \text{value_x} \rangle \Leftarrow \langle \text{value_1} \rangle$

$\langle \text{index_x} \rangle \Leftarrow \langle \text{index_1} \rangle$

Else

$\langle \text{value_x} \rangle \Leftarrow \langle \text{value_2} \rangle$

$\langle \text{index_x} \rangle \Leftarrow \langle \text{index_2} \rangle$

Equality is a bit cleverer – it rarely matters

If it does to you, see the MPI [standard](#)

Searching (2)

Operator `MPI_MINLOC` does precisely that

Operator `MPI_MAXLOC` searches for the maximum

You create the (`<value>`,`<index>`) pairs first

The `<index>` can be anything – whatever is useful

An `<index>` should usually be globally unique

I.e. **not** just the `index` into a `local array`

E.g. combine `process number` and `local index`

So how do we set up the data?

Fortran Searching (1)

Fortran 77 did not have derived types
The datatypes are arrays of length two

Recommended datatypes are:

MPI_2INTEGER & MPI_2DOUBLE_PRECISION

DOUBLE PRECISION can hold any INTEGER
on any current system, when using MPI

I don't recommend MPI_2REAL, except on Cray

Fortran Searching (2)

```
INTEGER :: sendbuf ( 2 , 100 ) ,      &  
          recvbuf ( 2 , 100 ) , myrank , error , i  
INTEGER, PARAMETER :: root = 3
```

```
DO i = 1 , 100  
    sendbuf ( 1 , i ) = <value>  
    sendbuf ( 2 , i ) = 1000 * myrank + i  
END DO
```

```
CALL MPI_Reduce ( sendbuf , recvbuf ,      &  
                  100, MPI_2INTEGER , MPI_MINLOC ,      &  
                  root , MPI_COMM_WORLD , error )
```

C Searching

The **datatypes** are “**struct {<value type>; int;}**”

- **C** structure **layout** is a can of worms

Recommended **datatypes** are:

MPI_2INT, **MPI_LONG_INT**, **MPI_DOUBLE_INT**

For **<value type>** of **int**, **long** and **double**

That will **usually** work – **not always**

Use **MPI_LONG_DOUBLE_INT** for “**long double**”

- Don't use **MPI_FLOAT_INT** or **SHORT_INT**
for **C** reasons you **don't** want to know!

C Example

```
struct { double value ; int index ; }  
    sendbuf [100] , recvbuf [100] ;  
int root = 3, myrank , error , i;  
  
for ( i = 1 ; i < 100 ; ++i ) {  
    sendbuf [ i ] . value = <value> ;  
    sendbuf [ i ] . index = 1000 * myrank + i ;  
}  
  
error = MPI_Reduce ( sendbuf , recvbuf ,  
    100, MPI_DOUBLE_INT , MPI_MINLOC ,  
    root , MPI_COMM_WORLD )
```

Data Distribution (1)

It can be inconvenient to make all **counts** the same
E.g. with a **100×100** matrix on **16** CPUs

One approach is to **pad** the short **vectors**

- That is usually more efficient than it looks

There are also extended MPI **collectives** for that
Obviously, their **interface** is more complicated

MPI_Gatherv, **MPI_Scatterv**,
MPI_Allgatherv, **MPI_Alltoallv**

- Use whichever approach is easiest for you

Data Distribution (2)

A **vector of counts** instead of a single **count**
One **count** for each **process**


- Provided only where there are **multiple buffers**
Receive counts for **MPI_Gatherv** & **MPI_Allgatherv**
Send counts for **MPI_Scatterv**
Both counts for **MPI_Alltoallv**

Used only on **root** for **MPI_Gatherv** & **MPI_Scatterv**



- But, for **MPI_Allgatherv** and **MPI_Alltoallv**,
the **count vectors** must match on all **processes**

I recommend **always** making them match, for sanity


Gatherv



Process 0




Process 1





Process 2





Allgatherv




Process 0



Process 1



Process 2



Data Distribution (3)

The **scalar counts** may all be different, of course
They must match for each **pairwise** send and receive

E.g. for **MPI_Gatherv**:

the **send count** on process **N** matches
the **Nth** receive count element on the **root**

MPI_Scatterv just the converse of **MPI_Gatherv**

For **MPI_Allgatherv**:

the **send count** on process **N** matches
Nth receive count element on **all** processes

Data Distribution (4)

The most complicated one is `MPI_Alltoallv`

It isn't hard to use, if you keep a clear head

Use a pencil and paper if you get confused

Consider processes `M` and `N`

the `Nth` send count on process `M`






matches the `Mth` receive count on process `N`

As said earlier, think of it as a `matrix transpose`!

With the `data vectors` as its `elements`

Alltoallv

Process 0



Process 2

Data Distribution (5)

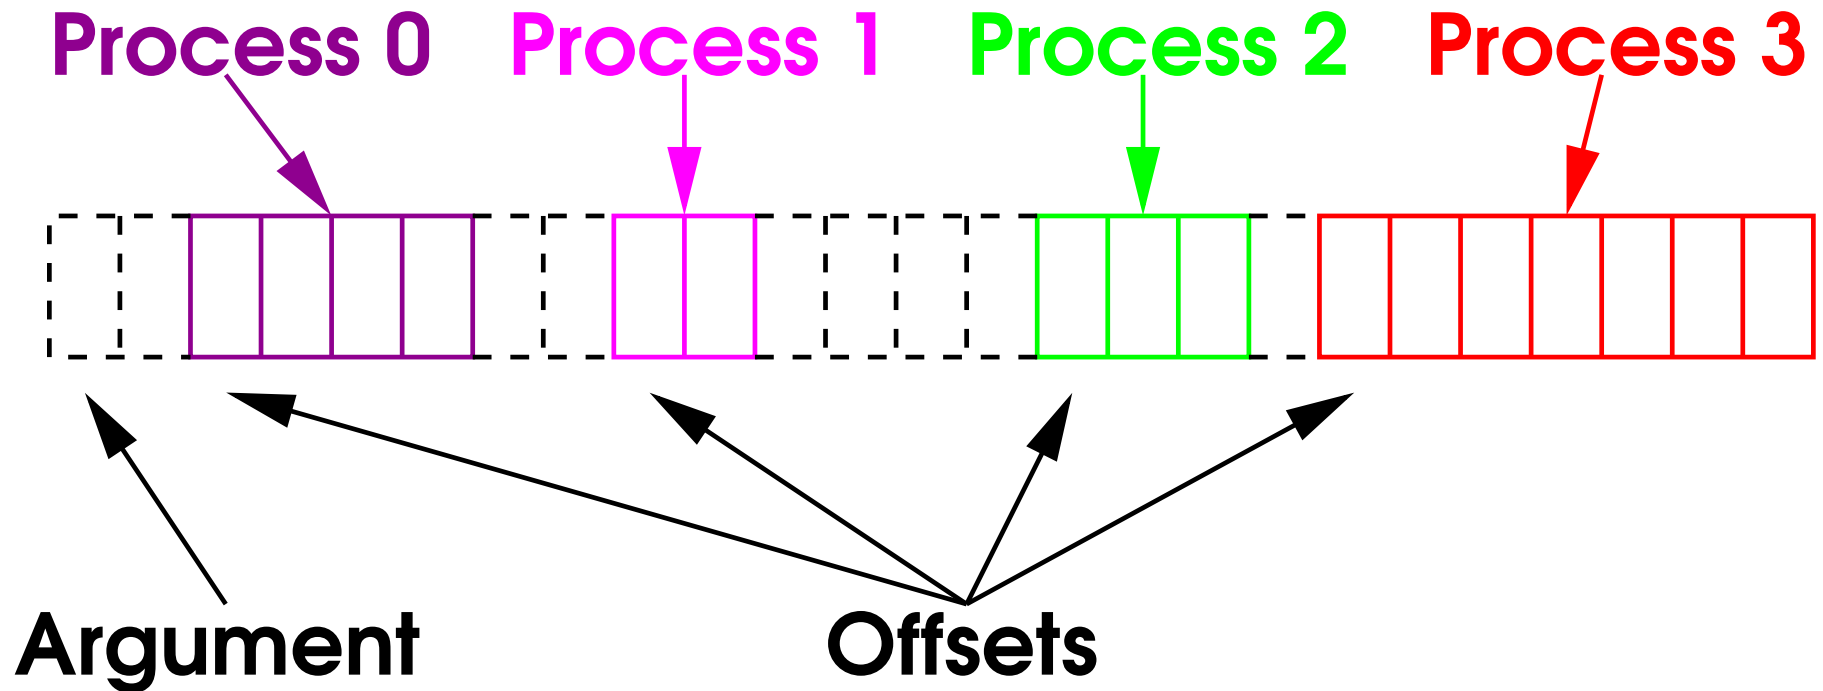
Where they have a **vector of counts**
they also have a **vector of offsets**

The **offset** of the data of each **process**
Not the **offset** of the basic **elements**

This allows for **discontiguous** buffers
Each **pairwise transfer** must still be **contiguous**

Normally, the first **offset** will be **zero**
But here is a picture of when it isn't

Multiple Transfer Buffers



Counts = (4 2 3 7)

Offsets = (2 8 14 18)

Data Distribution (6)

Unlike the **counts**, the **offsets** are purely **local**
They need not match on all **processes**

Even in the case of **MPI_Alltoallv**
the **offset vectors** needn't match in any way

Each one is used **just** as a mapping for
the **local** layout of its associated buffer

Data Distribution (7)





Keep your use of these **collectives** simple
MPI won't get confused, but you and I will
And any **overlap** is **undefined behaviour**

Will show a picture of a fairly general **MPI_Alltoallv**
Just to see what can be done, not to recommend it

Then simple examples of using **MPI_Gatherv**
Start with this (or **MPI_Scatterv**) when testing

Alltoallv

Process 0



Process 2

Gatherv (1)

Fortran example:

```
INTEGER , DIMENSION ( 0 : * ) :: counts
REAL(KIND=KIND(0.0D0)) ::      &
    sendbuf ( 100 ) , recvbuf ( 100 , 30 )
INTEGER :: myrank , error , i
INTEGER , PARAMETER :: root = 3 ,      &
    offsets ( * ) = ( / ( 100 * i , i = 0 , 30 - 1 ) / )

CALL MPI_Gatherv ( sendbuf , counts ( myrank ) ,      &
    MPI_DOUBLE_PRECISION ,      &
    recvbuf , counts , offsets ,      &
    MPI_DOUBLE_PRECISION ,      &
    root , MPI_COMM_WORLD , error )
```

Gatherv (2)

C example:

```
int counts [ ] ;  
double sendbuf[100] , recvbuf[30][100] ;  
int root = 3 , offsets[30] , error, i ;  
for ( i = 0 ; i < 30 ; ++ i )  
    offsets [ i ] = 100 * i ;  
  
error = MPI_Gatherv (  
    sendbuf , counts [ myrank ] , MPI_DOUBLE ,  
    recvbuf , counts , offsets , MPI_DOUBLE ,  
    root , MPI_COMM_WORLD ) ;
```

Practical Point

That's easy when the **counts** are **predictable**
But a lot of the time, they won't be

For **scatter**, calculate the **counts** on **root**
Use **MPI_Scatter** for the **count values**
Then do the full **MPI_Scatterv** on the data

Gather calculates a **count** on **each process**
Use **MPI_Gather** for the **count values**
Then do the full **MPI_Gatherv** on the data

Allgather and **alltoall** are similar

Efficiency (1)

- Generally, use **collectives** wherever possible
Provide most opportunity for **implementation** tuning
- Use the **composite** ones where that is useful
MPI_Allgather, **MPI_Allreduce**, **MPI_Alltoall**
- Do as much as possible in one **collective**
Fewer, **larger** transfers are always better

Consider packing **scalars** into **arrays** for **copying**
Even converting **integers** to **reals** to do so

Efficiency (2)

But what **ARE** 'small' and 'large'?

- A rule of thumb is about **4 KB** per buffer
That is each single, pairwise transfer buffer

But it's only a **rule of thumb** and not a hard fact

- Don't **waste time** packing data unless you **need to**

Synchronisation

Collectives are not synchronised at all
except for `MPI_Barrier`, that is

Up to **three** successive **collectives** can overlap
In theory, this allows for improved efficiency

In practice, it makes it hard to measure times

To synchronise, call `MPI_Barrier`

- The **first** process **leaves** only after
the **last** process **enters**

Scheduling (1)

Implementations usually tune for gang scheduling
Collectives often run faster when synchronised

- Consider adding a barrier before every collective
It's an absolutely trivial change, after all

Best to run 3–10 times with, and 3–10 without
The answer will either be clear, or it won't matter

If you have major, non-systematic differences
you have a nasty problem, and may need help

Scheduling (2)

You can overlap **collectives** and **point-to-point**
MPI requires **implementations** to make that work

- I **strongly** advise **not** doing that

A correct program will definitely not hang or crash
But it may run horribly slowly ...

Remember that **three** collectives can **overlap**?
Point-to-point can **interleave** with those

- I recommend **alternating** the modes of use
It's a lot easier to validate and debug, too!

Scheduling (3)

Start here ...

[Consider calling **MPI_Barrier** here]

Any number of **collective** calls

[Consider calling **MPI_Barrier** here]

Any number of **point-to-point** calls

Wait for **all** of those calls to finish

And repeat from the beginning ...

In-Place Collectives (1)

You can **usually** pass the same **array** twice
But it's a breach of the **Fortran** standard
And not **clearly** permitted in **C** †

- I recommend avoiding it if at all possible
It will rarely cause trouble – or get diagnosed
But, if it does, the bug will be almost unfindable

I have used systems on which it would fail
The **Hitachi SR2201** was one, for example

† Ask me why, offline, if you are masochistic

In-Place Collectives (2)

- **MPI 2** defined a **MPI_IN_PLACE** pseudo-buffer

Specifies the **result** overwrites the **input**

I.e. the real buffer is both **source** and **target**

Need to read MPI standard for full specification

For **allgather[v]**, **alltoall[v]** and **allreduce[v]**:

use it for the **send** buffer on **all** processes

Send **counts**, **datatype** and **offsets** ignored

In-Place Collectives (3)

Will give only a C example:

```
double buffer [ 30 ] [ 100 ] ;  
int error ;  
error = MPI_Alltoall (   
    MPI_IN_PLACE , 100 , MPI_DOUBLE ,  
/* Or even 'MPI_IN_PLACE , 0 , MPI_INT ,' */  
    buffer , 100 , MPI_DOUBLE ,  
    MPI_COMM_WORLD )
```

The Fortran is very similar

Epilogue

That is essentially **all** you need to know!

We have covered everything that seems to be used

MPI **collectives** look very complicated, but aren't

A few more features, which are rarely used

Mainly introduced by **MPI 2** for advanced uses

They are mentioned, briefly, in the extra lectures

Two exercises on **searching** and **scatterv**