

Programming with MPI

Error Handling

Nick Maclaren

nmm1@cam.ac.uk

May 2008

Error Handling (1)

Most **standards** get this hopelessly wrong
MPI gets it at least half right, as follows:

- All **invalid** uses are defined to be **erroneous**
Implementations are encouraged to detect them
- Most **undefined** results are detectable
E.g. they are set to a special, invalid value
- Most nonsense is defined to be **erroneous**
E.g. you cannot legally create a **deadlock**

Error Handling (2)

- No concept of conforming but undefined
I.e. it's a valid program with no known meaning

C is infested with it, Fortran/C++ use it
MPI has it, but only as bugs in the standard

- The default error handling is to stop
Not handling errors is fairly fail-safe in MPI

Error Handling (3)

However, all is not sweetness and light:

- Implementations not **required** to detect errors
Some errors are usually detected, others rarely are

- MPI has not specified a **debugging mode**
Error detection is at the whim of your implementor

Ones that can be detected **locally** usually are
e.g. providing an out-of-range **process number**
Inconsistencies across **collectives** **may** be

Error Handling (4)

- Some errors are almost undetectable
They include most language/MPI interface ones
E.g. incorrect datatype for the buffer type
 - Non-MPI ones are obviously not handled
And they may cause MPI to fail horribly
- MPI can't fix up other standards' defects!

Changing Error Handling

- There are programmable error handling facilities
But they don't allow actual **recovery**
Ask me if you want to know why this is unavoidable

You can use them **only** for cleaning-up
Including writing your own **diagnostics**

- But **do** look at the implementation documents
MPI encourages **documented** enhancements

Simple Error Handling (1)

There are several predefined **error handlers**

MPI_ERRORS_ARE_FATAL is the default

MPI_ERRORS_RETURN returns an **error code**

Fortran last argument, and **C** function result

Simple Error Handling (2)

- You attach the setting to a **communicator**
I recommend setting it **early**, and setting it **once**
And setting it **consistently** across **processes**

The call to do that is one that has changed name:

MPI_Comm_set_errhandler (new name)

MPI_Errhandler_set (old name)

If an MPI function returns an **error code**

i.e. anything that isn't **MPI_SUCCESS**

Call your code to diagnose, clean-up and stop

Simple Error Handling (3)

Error codes are implementation dependent

- A function to map them into an error string

You should use this for reasonable diagnostics

You do this by calling `MPI_Error_string`

Maps the error code to a textual message

Maximum length `MPI_MAX_ERROR_STRING`

- This is a block of text and not a C string

The length is returned via a separate argument

Warning – Edge of Cliff

- `MPI_ERRORS_RETURN` is dangerous

You **must** test for **errors** in **ALL** calls

One undetected **error** will cause **chaos** later

- But this facility can be very useful

You can write your own, **helpful** diagnostics

You can flush all your output to files

You can tidy up external state and not just crash

- Also can be used **temporarily** for debugging

Best to set the mode just around the failing call

Fortran Error Handling (1)

```
INTEGER :: error
```

```
CALL MPI_Comm_set_errhandler (    &  
    MPI_COMM_WORLD ,    &  
    MPI_ERRORS_RETURN , error )
```

Old versions of **gfortran** do not support this
There was a bug in **generic resolution** handling
There is a truly mind-boggling bypass if you need it

Fortran Error Handling (1.5)

This code works – heaven alone knows why!
It's harmless (but unhelpful) on other systems

```
INTEGER :: error, junk
```

```
junk = MPI_COMM_WORLD
```

```
CALL MPI_Comm_set_errhandler (      &  
    junk ,      &  
    MPI_ERRORS_RETURN , error )
```

Fortran Error Handling (2)

```
INTEGER :: error , length , temp
CHARACTER      &
    ( LEN = MPI_MAX_ERROR_STRING) :: message

< call some MPI function >
IF ( error /= MPI_SUCCESS ) THEN
    CALL MPI_Error_string ( error ,      &
        message , length , temp )
    PRINT * , message(1:length)
    CALL MPI_Abort ( MPI_COMM_WORLD ,      &
        1 , temp )
END IF
```

C Error Handling (1)

```
int error ;
```

```
error = MPI_Comm_set_errhandler (  
    MPI_COMM_WORLD ,  
    MPI_ERRORS_RETURN ) ;
```

C Error Handling (2)

```
int error , length ;  
char message [ MPI_MAX_ERROR_STRING + 1 ] ;  
  
< call some MPI function >  
if ( error != MPI_SUCCESS ) {  
    MPI_Error_string ( error , message ,  
        & length ) ;  
    procname [ length ] = '\0' ;  
    printf ("%s\n" , message) ;  
    MPI_Abort ( MPI_COMM_WORLD , 1 ) ;  
}
```

Note the way that the **length** is returned

Advanced Error Handling (1)

Can also map **error codes** into **error classes**
with the function **MPI_Error_class**

A **documented** set of **60** distinct values
with names **MPI_ERR_...**

Use these to distinguish various types of error

You may want to do this, for advanced handling
I can't offhand imagine why, but it's there

Advanced Error Handling (2)

Can define your own **error handlers**

Just functions to call when there is an **error**

- Safer than using **MPI_ERRORS_RETURN**

Provided that you code the function carefully

Just the same logic as for **MPI_ERRORS_RETURN**
in the examples that were given above

This course doesn't cover it, for simplicity

Experienced programmers will have no trouble

Advanced Error Handling (3)

Functions are:

`MPI_Comm_create_errhandler` (new name)

`MPI_Errhandler_create` (old name)

`MPI_Errhandler_free`

C type name `MPI_Handle_function`

And a MPI constant `MPI_ERRHANDLER_NULL`

Very Advanced Error Handling

- I recommend **not** doing any of this
Even most experts will never need or want to

Error handling is, in fact, purely **local**

Every **process** can have a different handler

Actually, every **communicator** in every **process** ...

Also, **MPI 2** extended it to some other **classes**

You can also change it whenever you want

For saving and restoring the old one, you need:

MPI_Comm_Get_errhandler (new name)

MPI_Errhandler_get (old name)

Epilogue

That's more-or-less all there is to MPI errors

The only exercise is to try the simple case out