

Programming with MPI

One-sided Communication

Nick Maclaren

nmm1@cam.ac.uk

October 2010

What Is It?

This corresponds to what is often called **RDMA**
That is “**Remote Direct Memory Access**”

[It is called **RMA** in MPI]

Much loved by **Cray** and the **DoD/ASCI** people

One **process** accesses the **data** of another

Potentially more efficient on **SMP** systems


It was added in **MPI 2**, but is not often used

- I will explain why I don't like **RDMA**
Then describe how it can be used **semi-safely**

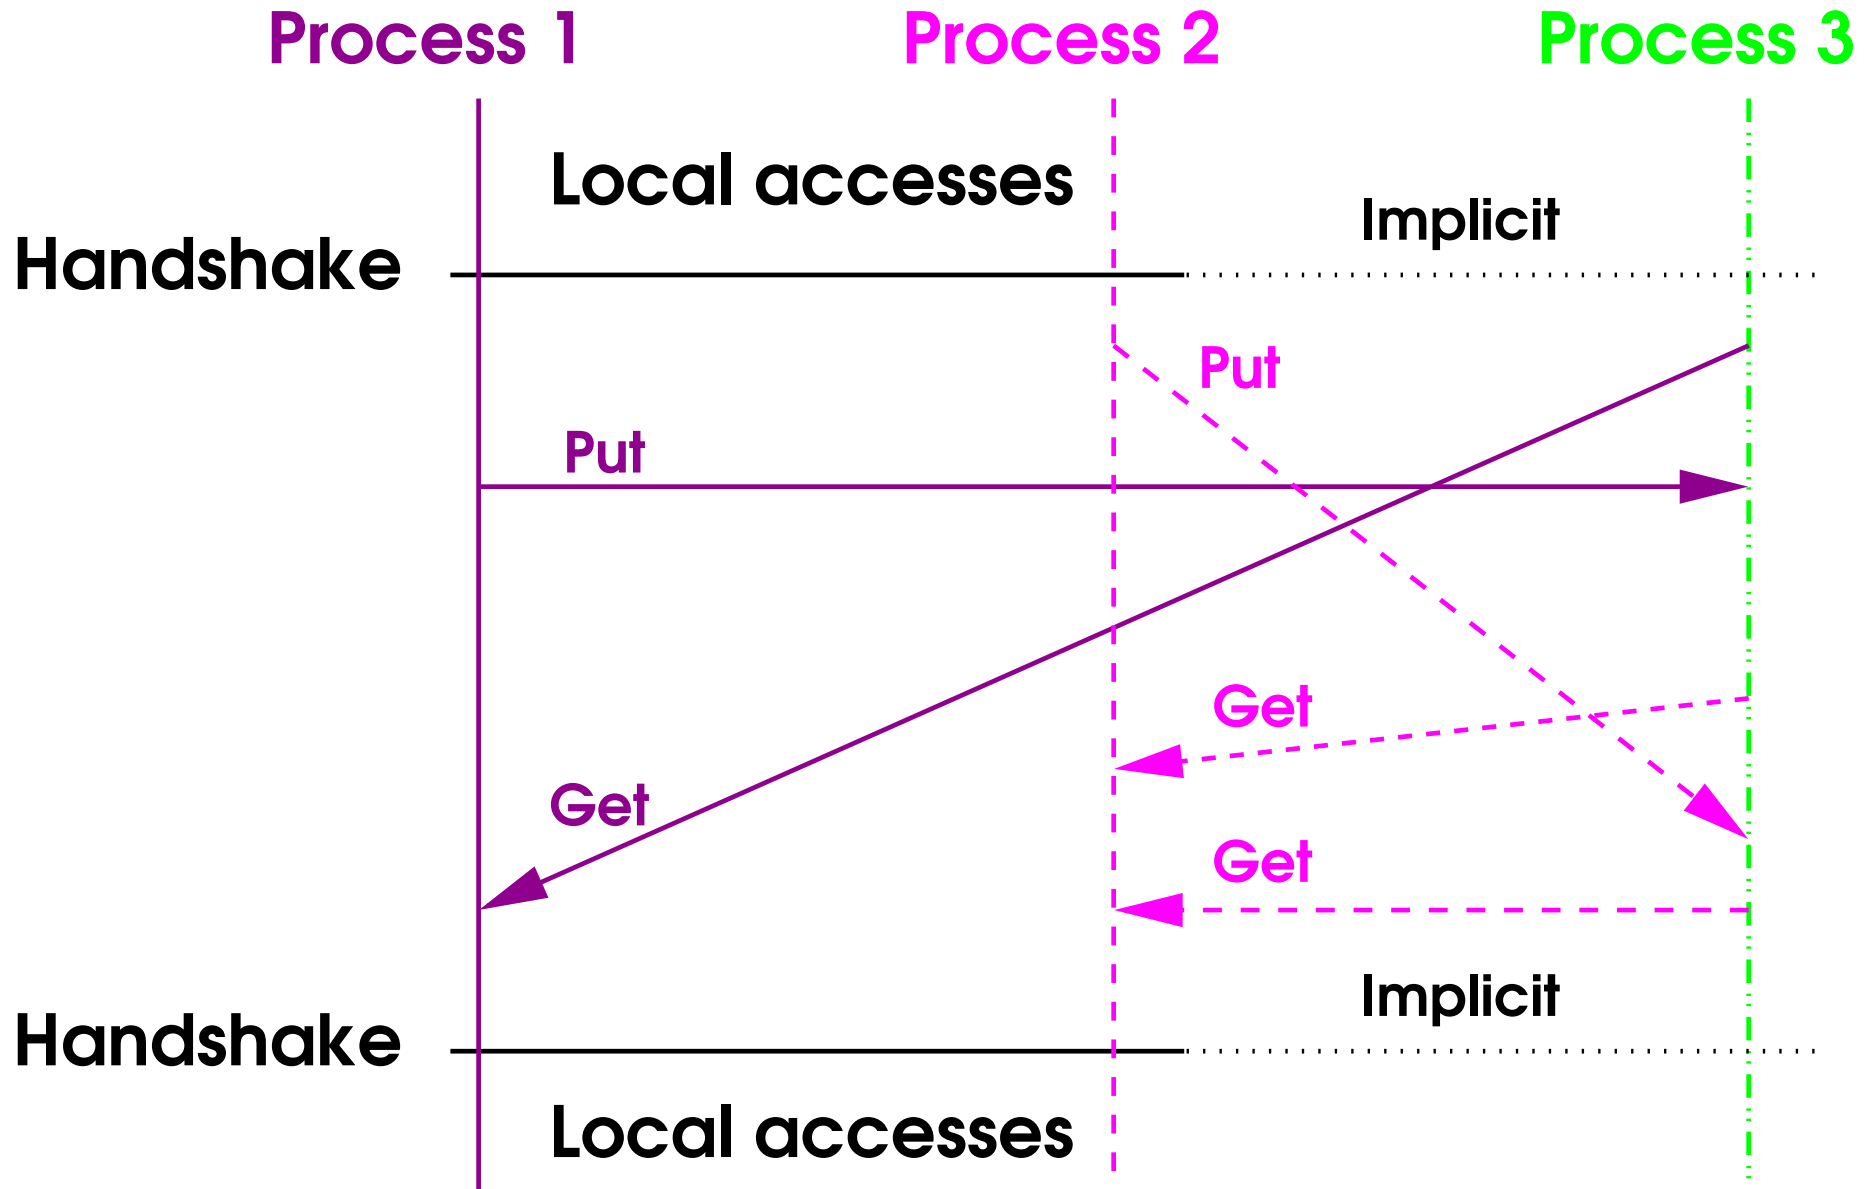
RDMA Models

- Active RDMA is **batching up** sends and receives
A bit like **non-blocking** using **MPI_Waitall**
Not too hard to use correctly, but needs **discipline**
- Passive uses **uninvolved** process's memory
Much harder to use and **less portable**
PGAS languages effectively use this model
- Lastly, can use true **virtual shared memory**
Using **all processes'** memory as if it were **local**
MPI does not support this, for very good reasons

Active (Fenced) Communication



Passive Communication



Problems with RDMA (1)

RDMA is often said to be **easier** to use **correctly**

Twaddle!

How does **other process** know its **data** is in use?
Get it wrong, and almost **unfindable chaos** ensues

- Adding **handshaking** often makes it pointless

Most **specifications** don't say how to handle that
MPI does, which makes it fairly complicated

Problems with RDMA (2)

Another problem is that it handicaps **analysis tools**
With **two-sided**, can diagnose unmatched **transfers**
But there's no concept of matching for **one-sided**

Its **semantics** are similar to **SMP threading**

- Lots of lovely potential **race conditions**
- And no **realistic tools** to detect them

Similar remarks apply to non-trivial **tuning**

Those problems are **insoluble** – provably so
Most people's experiences are very negative

Problems with RDMA (3)

Next problem is **implementation** and **progress**
Passive RDMA needs a **remote agent** on memory

May hang until **far process** reaches an **MPI call**
With **passive** use, that **may never happen**
Can also occur with **active**, but **less likely**

- **Deadlock** may happen even in **correct code**
MPI forbids that for **two-sided** and (mostly) **active**

Some **subtle problems** even on **SMP** systems

Problems with RDMA (4)

Last problem is **performance** complications

- On **commodity clusters**, it may well be **slower**
Remote agent may interfere with **MPI process**

- Probably **faster** on specialist **HPC clusters**
Cray etc. have **hardware and kernel** RDMA
Transfers data without any action by MPI process

- On **SMP** systems, it's very **hard to guess**
 - + May well have **lower communication** overhead
 - But may cause serious **cache** and **TLB** problems

Most Likely Benefit

- One **logical** transfer is many **actual** ones
I.e. you need to transfer **lots of objects** together
Essentially an alternative to **packing up** into one array
- On **SMP** system or specialist **HPC** cluster
Does **NOT** gain you anything using **Ethernet**
Nor **Infiniband** with **commodity** software

TCP/IP and **Ethernet** need a **kernel** thread
Currently, so does the **OpenIB** Infiniband driver

A SMP System Myth

The following is **NOT** true, in general:

One-sided transfers use fewer **memory copies**

True in **theory**, but generally **not** in **practice**

Probably true only for **MPI_Alloc_mem** memory

- For all **Unix-like** systems, including **Microsoft's**
Process **A** cannot access process **B's** memory

- Can only if using a **shared memory segment**
And almost all data areas **aren't** in one of them
⇒ Two memory copies needed for it, as well

Recommendation

- Stick to simple forms of **active** communication
Only one taught here is **collective-like** transfers
Group pairwise transfers are also described, roughly
 - Be extremely cautious about **deadlock**
Use only methods that **don't rely on** progress details
 - Scrupulously adhere to **MPI's restrictions**
Designed to maximise **reliable** implementability
- That is, naturally, all that **this course** teaches

Preliminaries

Recommended to create **info object**, for **efficiency**

Most of its uses are for other **MPI 2** facilities

Do this on **each process**, but it is **not collective**

- Asserts that won't use **passive** communication
Curiously, no option to specify a **read-only window**

Also a relevant function **MPI_Alloc_mem**

May help with performance on **some systems**

But is **NOT** recommended for use in **Fortran**

It's fine in **MPI 3.0**, which uses **Fortran 2003**

Fortran Example

```
INTEGER :: info , error
```

```
CALL MPI_Info_create ( info , error )
```

```
CALL MPI_Info_set ( info , "no_locks" ,      &  
    "true" , error )
```

! Use the info object

```
CALL MPI_Info_free ( info , error )
```

C Example

```
MPI_Info info ;  
int error ;  
  
error = MPI_Info_create ( & info ) ;  
error = MPI_Info_set ( info , "no_locks" , "true" ) ;  
  
/* Use the info object */  
  
error = MPI_Info_free ( info ) ;
```

Windows

A **window** is an array used for one-sided **transfers**

- Must **register** and **release** such arrays **collectively**

Args are **address**, **size** in **bytes** and **element size**

- All of them can be **different** on each **process**

Same remark applies to the **element type** of the array

- **Offsets** are specified in units of **element size**

Like **array indices** into the **target window**

- **Size** of zero means no local memory **accessible**

Alignment

- Strongly advise to **align** correctly for **datatype**
Only a **performance** issue, but may be significant

Some systems may benefit from **stronger alignment**
Read the **implementation notes** to find that out

- **Element size** need not match the **element type**
But **simplest** use is to **make it match**

It can be **1** if using **byte offsets**

- In that case, it's **your** job to get it right

Fortran Example

```
REAL ( KIND = KIND ( 0.0D0 ) ) :: array ( 1000 )  
INTEGER :: info , window , size , error  
INTEGER ( KIND = MPI_ADDRESS_KIND ) :: win_size
```

! Create the info object as described above

```
CALL MPI_Sizeof ( array , size , error )
```

```
win_size = 1000 * size
```

```
CALL MPI_Win_create ( array , win_size ,      &  
                    size , info , MPI_COMM_WORLD ,      &  
                    window , error )
```

```
CALL MPI_Win_free ( window , error )
```

C Example

```
double array [ 1000 ] ;
MPI_Info info ;
MPI_Win window ;
int error ;
MPI_Aint win_size ;

/* Create the info object as described above */
win_size = 1000 * sizeof ( double ) ;
error = MPI_Win_create ( array , win_size ,
    sizeof ( double ) , info , MPI_COMM_WORLD ,
    & window ) ;

error = MPI_Win_free ( window ) ;
```

Horrible Gotcha

- The window must be an **ordinary** data array

Not an MPI standard issue, and is **not stated** there

It's a **hardware**, **system** and **compiler** one

Applies to **other** RDMA and **asynchronous** use, too

- Do not use **Fortran** **PARAMETER** arrays
- Do not use **C/C++** **static const** arrays
- Or anything **any library** returns that **might be**
- Or anything else that might be **exceptional**

At **best**, the **performance** might be **dire**

Fencing

Exactly like `MPI_Barrier`, but on a `window` set
`Assertions` described shortly – ignore for now

Fortran example:

```
CALL MPI_Win_fence ( assertions , window , error )
```

C example:

```
error = MPI_Win_fence ( assertions , window )
```

C example:

```
window.Fence ( assertions )
```

Use of Windows (1)

Rules for **use of windows** apply much more **generally**
But are **easiest to describe** in terms of fencing

Fences divide **time** into sequence of **access epochs**
Each **window** should be considered **separately**

⇒ So consider **one epoch** on **one window**

No restrictions if **not accessed remotely** in epoch

- It includes **local writes from it** using RDMA

That is **unlike** the rules for **MPI_Send**

Use of Windows (2)

Window is **exposed** if it may be **accessed remotely**
If the window is **exposed**:

- Mustn't **update any part of window** locally
Seems unnecessary, but are **good reasons** for this
Can use RDMA to **local process** to bypass this

- Any **location** may be **updated only once**
And **not at all** if it is **read** locally or remotely
The standard **write-once-or-read-many** model

Accumulations are an exception – see later

Optimisation Issues

Bends language standards in several subtle ways

- Can cause accesses to get out of order

Very rare for C and C++ – and truly evil if it does

- Almost all such problems are user error

Fortran is more optimisable, and it can happen

- Window should have ASYNCHRONOUS attribute

So should any dummy argument that it is passed to during any access epoch

- Simplest not to pass it during an access epoch

Assertions (1)

This is an **integer** passed to **synchronisation** calls
Can be combined using **logical OR** (**IOR** or '|')
Purely optional, but may help with **performance**

- If you get them **wrong**, behaviour is **undefined**
Pass the argument as **0** if you are unsure
- Fall into two classes: **local** ones and **collective**
Latter are supported **only** by **MPI_Win_fence**
- Apply to an **epoch** between **synchronisations**
Can be either **preceding** or **succeeding** epoch

Assertions (2)

Local assertions:

MPI_MODE_NOSTORE – about preceding epoch

The window was not updated in any way

MPI_MODE_NOPUT – about succeeding epoch

The window will not be updated by RDMA

Collective assertions:

MPI_MODE_NOPRECEDE – about preceding epoch

No RDMA calls were made by this process

MPI_MODE_NOSUCCEED – succeeding epoch

No RDMA calls will be made by this process

Assertions (3)

- If **alternating** computation and communication
Do the following **collectively** (same on **all processes**)

- . . . Do some **computation**
- **Fence** with **MPI_MODE_NOPRECEDE**
- Do some RDMA **communication**
- **Fence** with
MPI_MODE_NOSUCCEED | MPI_MODE_NOPUT
- Do some **computation** . . .

And use **MPI_MODE_NOSTORE** when appropriate
It does not need to be the same on all processes

Transfers

`MPI_Put` and `MPI_Get` have identical syntax

Effectively start a remote `MPI_Recv` or `MPI_Send`

- Both sets of arguments are provided by the caller
Remote ones interpreted in context of target window
- Strongly recommended to match types and counts
Remote datatype must match remote element type

Like non-blocking, so don't reuse local buffers

The next synchronisation call completes the transfer

- No guarantee that it is implemented like that

Fortran Example

```
REAL :: array_1 ( 1000 ) , array_2 ( 1000 )  
INTEGER :: to = 3 , from = 2 , window, error
```

```
CALL MPI_Put ( array_1 , 1000 , MPI_REAL ,      &  
              to , offset_1 , 1000 , MPI_REAL ,      &  
              window , error )
```

```
CALL MPI_Get ( array_2 , 1000 , MPI_REAL ,      &  
              from , offset_2 , 1000 , MPI_REAL ,      &  
              window , error )
```

C Example

```
double array_1 [ 1000 ] , array_2 [ 1000 ] ;
```

```
MPI_Win window ;
```

```
int to = 3 , from = 2 , error ;
```

```
error = MPI_Put ( array_1 , 1000 , MPI_DOUBLE ,  
                to , offset_1 , 1000 * size , MPI_DOUBLE ,  
                window )
```

```
error = MPI_Get ( array_2 , 1000 , MPI_DOUBLE ,  
                from , offset_2 , 1000 * size , MPI_DOUBLE ,  
                window )
```

Reminder

- The **target address** is not specified directly

window describes the **remote array** to use

offset is in units of **remote element size**
from start of **remote window**

Accumulation (1)

- A point-to-point reduction with a remote result
Accumulation is done in an undefined order

Syntax and use is almost identical to `MPI_Put`

Reduction operation before the window argument

Exactly the same operations as in `MPI_Reduce`

Only predefined operations – not user-defined

One extra predefined operation `MPI_REPLACE`

Simply stores the value in target location

Accumulation (2)

- You mustn't **update or access** it any other way

But you can use multiple **separate** locations

Each can use **the same** or **a different** operation

- **Cannot** use it for **atomic** access

Can only **accumulate** until next **synchronisation**

Partly possible, but needs features not taught here

Use only **one operation** for one **location**

MPI doesn't require it – your sanity does

Syntax Comparison

```
CALL MPI_Put ( array , 1000 ,      &  
             MPI_REAL , to , offset , 1000 ,      &  
             MPI_REAL , window , error )
```

```
CALL MPI_Accumulate ( array , 1000 ,      &  
                   MPI_REAL , to , offset , 1000 ,      &  
                   MPI_REAL , MPI_SUM , window , error )
```

Note only change is addition of **MPI_SUM**

C and **C+** have identical changes

Writing Collectives

Above is enough to write **collectives** using RDMA
If **clean and simple**, you should have no trouble

MPI specifies to **never deadlock** but, for **sanity**:

- Do not overlap use of **window sets**
Like communicators, **serialise overlapping** use
- Do not overlap with other **types of transfer**
Serialise **collectives**, **point-to-point** and **one-sided**

Group Pairwise RDMA (1)

- This is a **lot** more **complicated**
So much so, this course **doesn't cover** the details
- It is also a **lot** more **error-prone**
Make a **mistake**, and your program may **deadlock**
Or it may cause **data corruption**, with no diagnostic
- Worse, these may happen only **probabilistically**

I **do not recommend** using this, except for experts
The following is a **summary** of how to use it

Group Pairwise RDMA (2)

`MPI_Win_post (group , assertions , window)`

This **registers** the window for RDMA **accesses**
Only the processes in **group** may **access** it
MPI specifies that it **does not block**

`MPI_Win_wait (window)`

This **blocks** until all RDMA has completed
All processes have called **MPI_Win_complete**
It **cancels** the **registration** for RDMA accesses

Group Pairwise RDMA (3)

`MPI_Win_start (group , assertions , window)`

This **registers** the window for RDMA **calls**

You may **access** only the data of processes in **group**

It may **block**, but is **not required** to

I.e. until all processes have called `MPI_Win_post`

`MPI_Win_complete (window)`

This **cancels** the **registration** for RDMA calls

It may **block**, but is **not required** to

Circumstances too complicated to describe here

Group Pairwise RDMA (4)

You will need some of the **group** calls

You **need not** use them **collectively** for this

See the lecture on **Communicators etc.** for them

MPI_Comm_group (comm , group)

This obtains the **group** for a **communicator**

MPI_Group_incl (group , count , ranks , newgroup)

This creates a **subset group** from some **ranks**

MPI_Group_free (group)

This frees a **group** after it has been used

Group Pairwise RDMA (5)

That's essentially **all of the facilities**

You can **test** for completion using **MPI_Win_test**
But **can't touch** the window until it **completes**

If you want to use group pairwise RDMA:

Read the MPI standard

Passive RDMA

Did you think that the **group pairwise** form is bad?
Passive RDMA is **simpler** but **much** trickier

Don't go there

Debugging and Tuning

You're on your own