

# Programming with MPI

## *Advanced Completion Issues*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

May 2008

# More on Completion

More on Point-to-Point made simplifying assumptions

This describes when those are **not** so

Three more advanced features complicate things

- Waiting for a **subset** of requests (described shortly)
- **Cancellation** of requests (described shortly)
- **Persistent** requests (described in a later lecture)

I don't recommend using any of these

But this is a description of the issues

# Empty Statuses

MPI has the concept of an **empty status**

An **empty status** looks like the following:

- The **tag** is **MPI\_ANY\_TAG**
- The **source** is **MPI\_ANY\_SOURCE**  
[ **⇒** or possibly **MPI\_PROC\_NULL** ]
- **MPI\_Get\_count** returns zero

And, for properties we haven't covered yet:

- The **error code** is **MPI\_SUCCESS**
- **MPI\_Test\_cancelled** returns **False**

# Completion of Subsets (1)

`MPI_Testsome` and `MPI_Waitsome`

These check for or complete **some** of the **requests**  
and return a **count** of how many  
plus arrays of **indices** and **statuses**

For **wait** and when **test**'s flag is **True**:

The **index array** lists the completed **requests**  
First **count** elements of the **status array** are set  
The other **statuses** are not defined

## Completion of Subsets (2)

If not enough of the **requests** are ready

- The **tests** set their **flag** to **False**
- The **waits** hang until something happens

If enough of the **requests** are ready

- **Any** completes just one **request**
- **Some/all** complete all ready **requests**
- The **tests** set their **flag** to **True**

All **completed requests** are released exactly  
as for the **individual request** forms

# Error Codes (1)

What if **not** using **MPI\_ERRORS\_ARE\_FATAL**?

Multiple **errors** from the **all** and **some** forms

One of the many reasons the default is easiest

The **error code** may be **MPI\_ERR\_IN\_STATUS**

The individual **error codes** are in the **statuses**

Including the **empty statuses** of the **all** forms

## Error Codes (2)

<status array> ( MPI\_ERROR , <index> ) (Fortran)

<status array> [ <index> ] . MPI\_ERROR (C)

The MPI\_ERROR fields are set if and only if:

- You call one of the **all** or **some** forms
- Its **error code** is MPI\_ERR\_IN\_STATUS

That field is never set for the **any** forms

I.e. exactly like the **individual request** forms

They will never return MPI\_ERR\_IN\_STATUS

# Fortran Multiple Errors

```
INTEGER :: i , error , requests ( 100 ) ,      &  
        statuses ( MPI_STATUS_SIZE , 100 )  
  
CALL MPI_Waitall ( 100 , requests , statuses , error )  
IF ( error == MPI_ERR_IN_STATUS ) THEN  
    DO i = 1 , 100  
        IF ( statuses ( MPI_ERROR , i ) /=      &  
            MPI_SUCCESS ) THEN  
            CALL fail ( statuses ( MPI_ERROR , i ) )  
        END IF  
    END DO  
ELSE IF ( error /= MPI_SUCCESS ) THEN  
    CALL fail ( error )  
END IF
```



# C Multiple Errors

```
int i , error , requests [ 100 ] ;  
MPI_Status statuses [ 100 ] ;  
  
error = MPI_Waitall ( 100 , requests , statuses ) ;  
if ( error == MPI_ERR_IN_STATUS ) {  
    for ( i = 1 ; i < 100 ; ++i ) {  
        if ( statuses[i] . MPI_ERROR !=  
            MPI_SUCCESS )  
            fail ( statuses[i] . MPI_ERROR )  
    }  
else if ( error != MPI_SUCCESS )  
    fail ( error ) ;
```

# Completion Oddities (1)

There are actually some exceptions to the above

- You can avoid them by not causing them

The facilities are described (briefly) later

- Persistent requests are not released  
you have to release them yourself

This course doesn't describe these in detail

- Cancellation is different from completion  
the request merely becomes inactive

You still have to complete or release it

# Completion Oddities (2)

Requests become **inactive** in only three ways:

1. Setting **MPI\_REQUEST\_NULL** explicitly
2. Passing an already **completed** request
3. Using **cancellation** (see later)

**Multiple completion** unavoidably causes 2

- Either remove them from the **request array**
- Or you can learn more about the functions

It isn't hard, but each group is different

# Inactive Requests (1)

We first consider the **individual request** forms

**Wait** and **test** work on **inactive requests**

- they return immediately and successfully
- the **status** is set to **empty**

# Inactive Requests (2)

We next consider the **any** forms

If none of the **requests** are **active**

Including the case of a zero length **request** array

- they return successfully and immediately
- the **index** is set to **MPI\_UNDEFINED**
- the **status** is set to **empty**

Otherwise, they consider just the **active requests**

I.e. very like the **individual request** forms

# Inactive Requests (3)

We now consider the **all** forms

If none of the **requests** are **active**

Including the case of a zero length **request** array

- they return successfully and immediately

Otherwise, they consider just the **active requests**

- In both cases, all **statuses** corresponding to **inactive requests** are set to **empty**

# Inactive Requests (4)

We last consider the **some** forms

If none of the **requests** are **active**

Including the case of a zero length **request** array

- they return successfully and immediately
- the **index count** is set to **MPI\_UNDEFINED**

Otherwise, they consider just the **active requests**

- The **index array** is only **completed requests**  
i.e. ones **completed** by this call
- Only completed **requests** have **statuses**

# Inactive Requests (5)

The above all looks like unnecessary complexity

- But it isn't – MPI has got it right

It means that you can write clean, obvious code  
And everything will all work as it should



# Cancellation (1)

This is just an overview of the facility

You may need to abandon active **requests**

⇒ Try to avoid ever getting into that hole

- Cancellation is for exceptional circumstances  
It may be both **unreliable** and **inefficient**

**MPI\_Cancel** will **start** the cancellation

- It will **not** release the **request**

## Cancellation (2)

You must still call `MPI_Wait` or `MPI_Test`  
Or one of the `request array` versions of those

- `MPI_Test_cancelled` checks the `status`  
Returns a flag saying if the cancellation succeeded
- If you use cancellation, test that `first`  
All other `status` fields are `undefined` if cancelled

# Cancellation (3)

You can also simply release the **request**  
By calling **MPI\_Request\_free**

You can also call this on active **requests**  
They will be disconnected, but will complete

- **DON'T** do that – not even for **sends**

You have no way of telling when they complete  
And what happened when they finally do